



上海人工智能实验室  
Shanghai Artificial Intelligence Laboratory



SCHOOL OF  
COMPUTING &  
DATA SCIENCE  
The University of Hong Kong

# Code Intelligence

Qiushi Sun

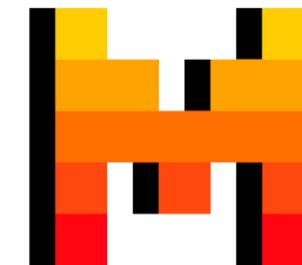
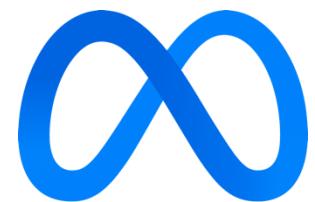
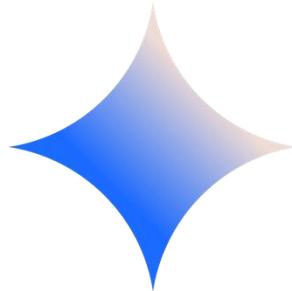
[qiushisun.github.io](https://qiushisun.github.io)

✉ @qiushi\_sun

# The Rise and Potential of Neural Code Intelligence

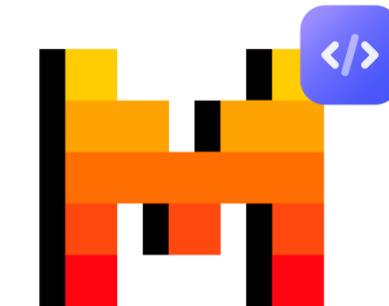
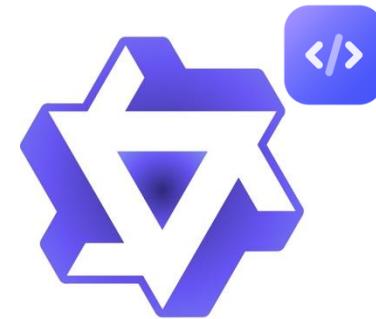
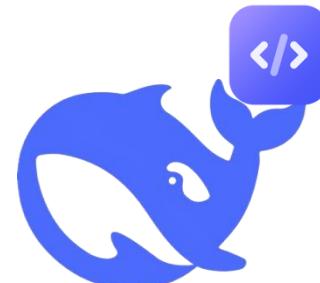
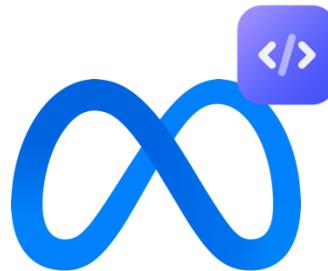
# Large Language Models

We are quite familiar with them



# Large Language Models: for Code

Code variants of LLMs





上海人工智能实验室  
Shanghai Artificial Intelligence Laboratory



National University  
of Singapore



香港大學自然語言處理實驗室  
Natural Language Processing Group, The University of Hong Kong



NANYANG  
TECHNOLOGICAL  
UNIVERSITY  
SINGAPORE



SCHOOL OF DATA  
SCIENCE & ENGINEERING  
数据科学与工程学院



# A Survey of Neural Code Intelligence: Paradigms, Advances and Beyond

Qiushi Sun, Zhirui Chen, Fangzhi Xu, Chang Ma, Kanzhi Cheng,  
Zhangyue Yin, Jianing Wang, Chengcheng Han, Renyu Zhu,  
Shuai Yuan, Pengcheng Yin, Qipeng Guo, Xipeng Qiu, Xiaoli Li,  
Fei Yuan, Lingpeng Kong, Xiang Li, Zhiyong Wu

[arXiv: 2403.14734 \[v5\] Sun, 26 Jan 2025](https://arxiv.org/abs/2403.14734)

# The Develop Timeline of CodeLMs

1. An increasing number of researchers are diving into
2. It is generally positively correlated with the development of LMs.

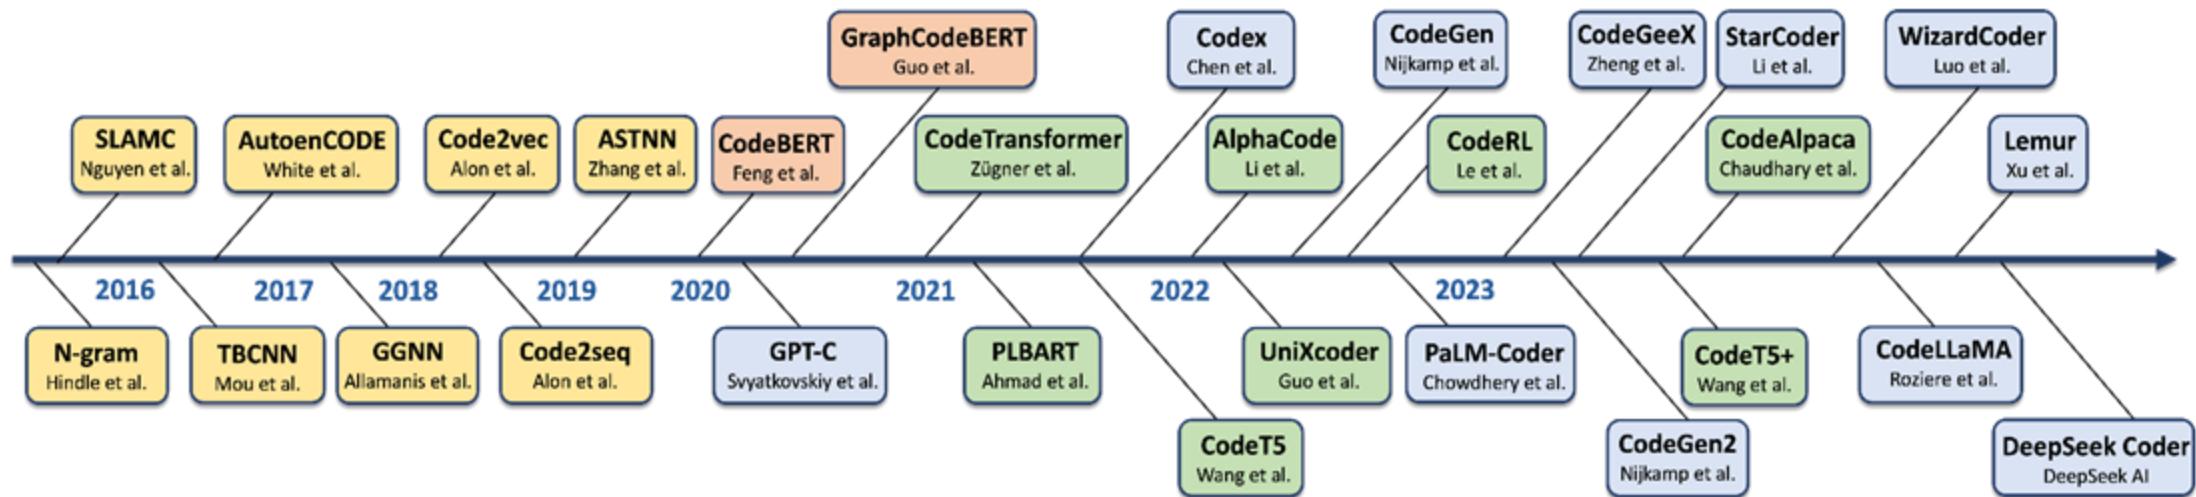
Papers are usually pubed at:

1. ML venues: NIPS, ICLR, ICML ...
2. NLP venues: \*ACL, COLM, ...
3. SE venues: ICSE, ASE, ISSTA, ...



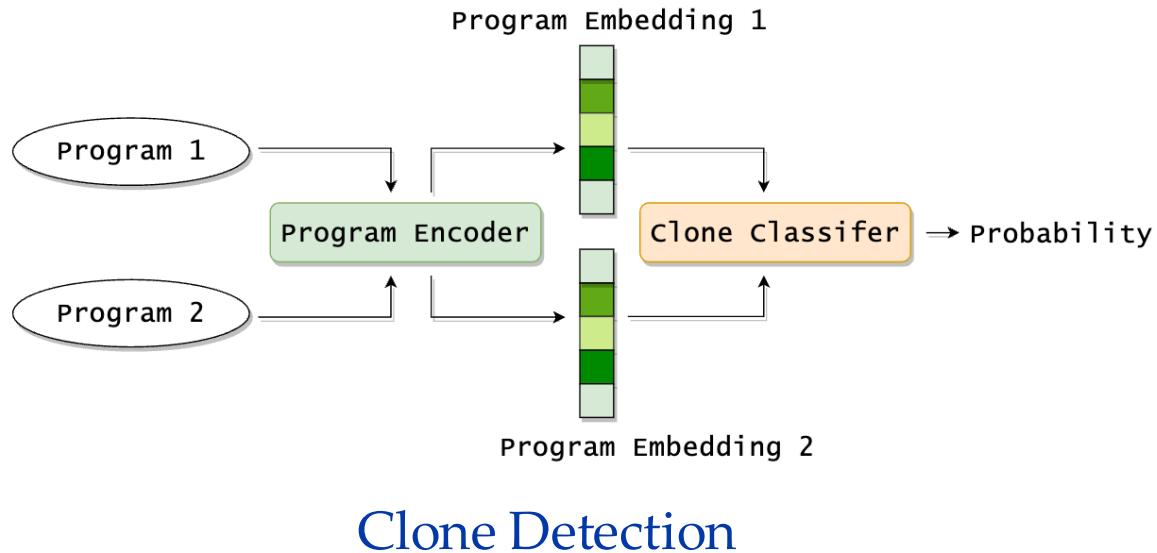
# The Develop Timeline of CodeLMs

A story through models, but today, we're focusing on code data.

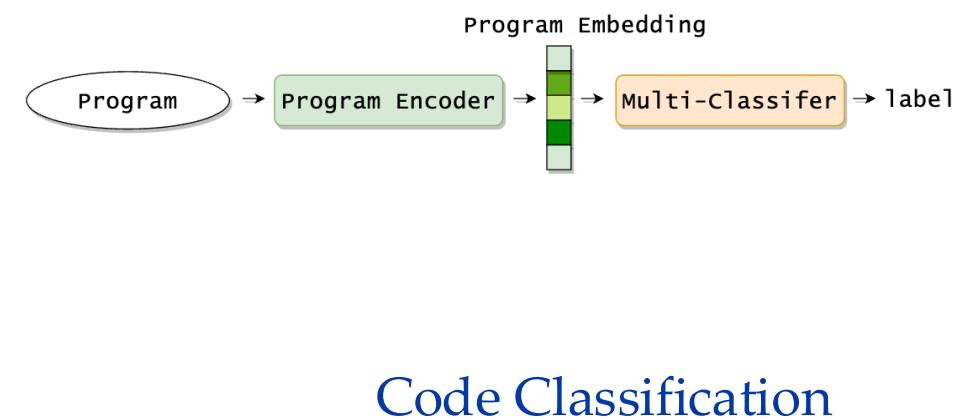


# Code-Related Tasks

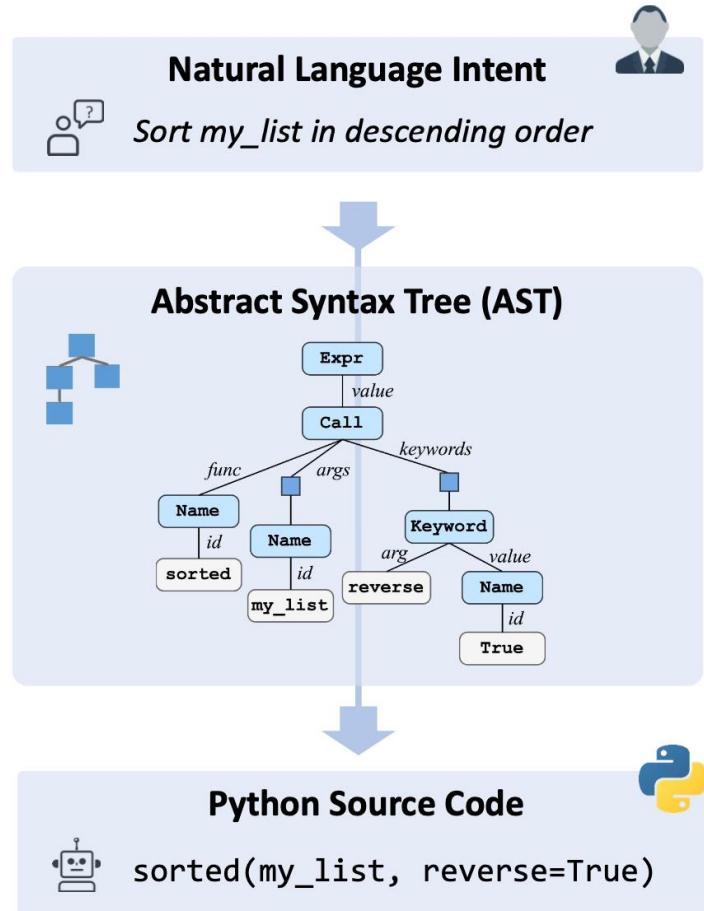
How it starts?



Before LLMs, we were most focused on how to construct code representations.



# Code Representation Learning

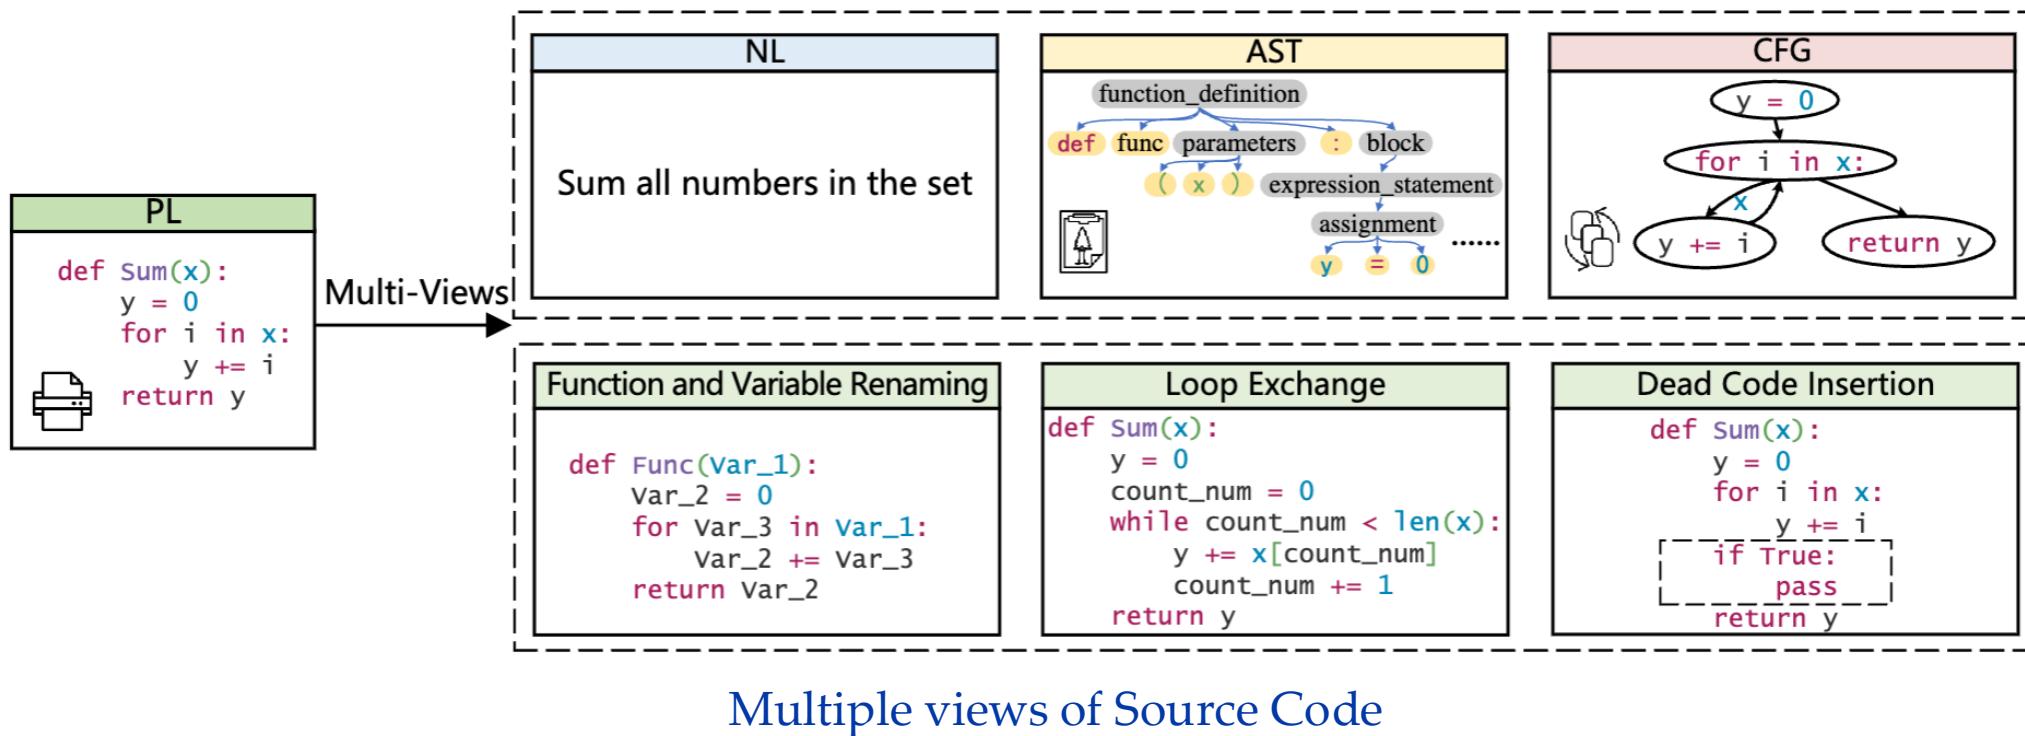


- Use Abstract Syntax Trees as general-purpose intermediate meaning representations
- $p_\theta(\text{AST} | \text{User Input})$  is a seq-to-tree model using program grammar as prior syntactic knowledge to constrain decoding space
- Deterministic transformation to source code

The pastoral era of code generation and understanding

# Code-Related Tasks

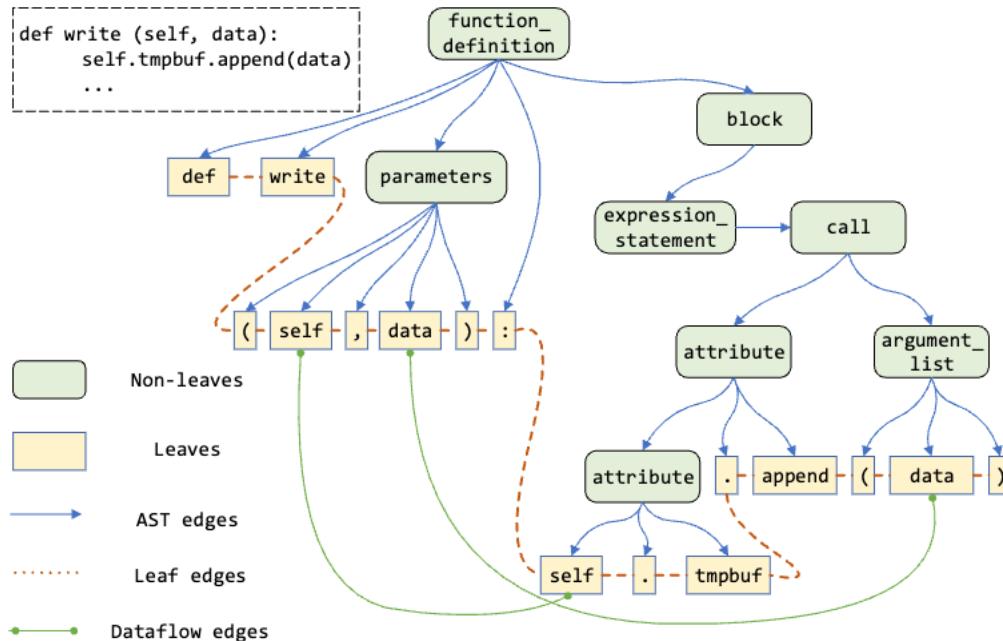
## How code differ from NL



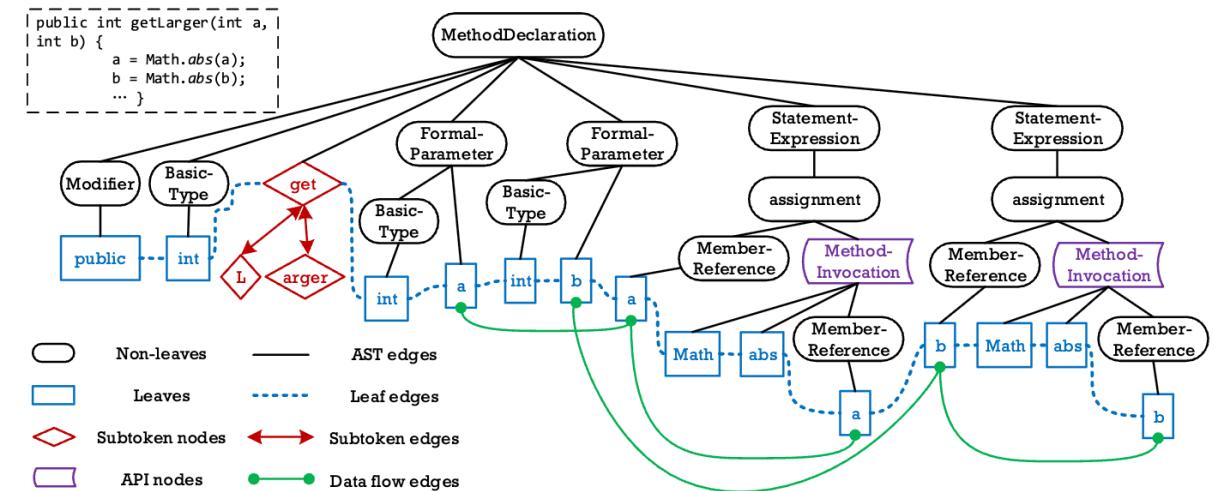
# Solving Code-Related Tasks

Section 2.1

How code differ from NL



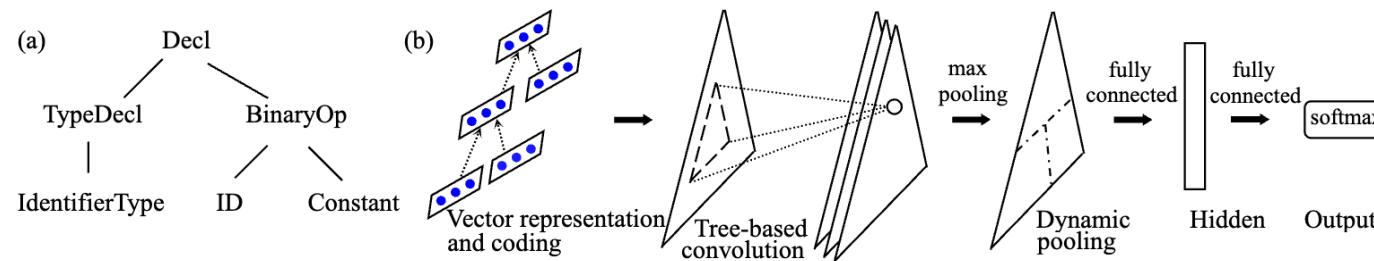
“Enhanced” ASTs



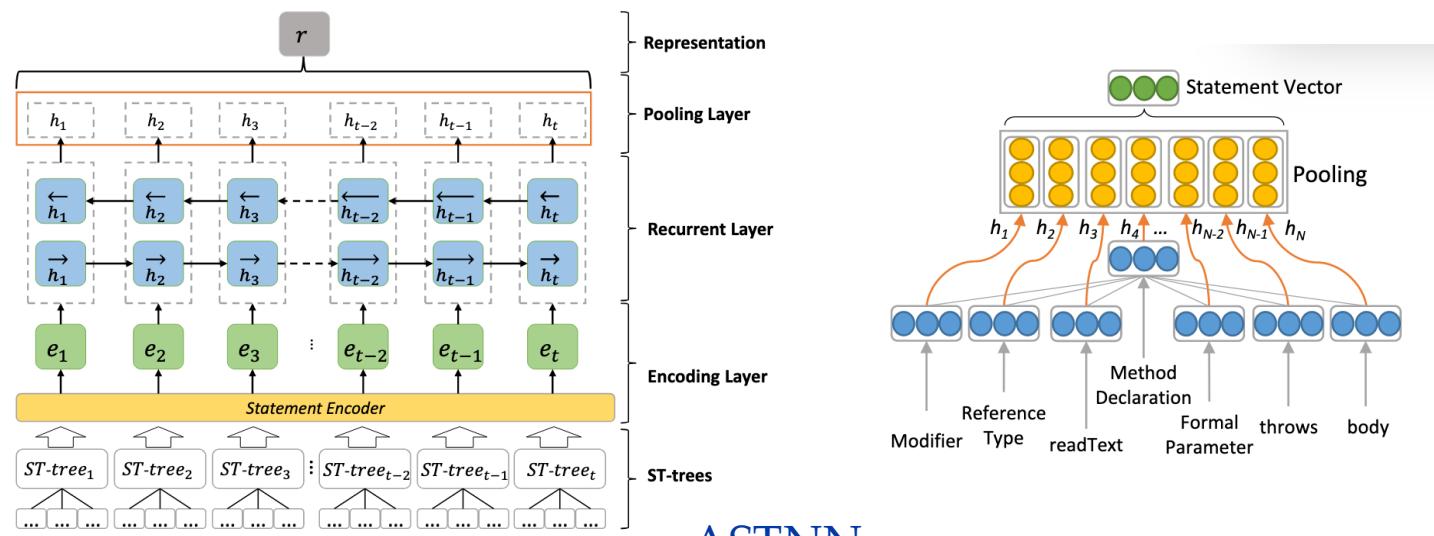
Another “Enhanced” ASTs

# Typical CodeLMs before Transformer

Section 2.2



TBCNN



ASTNN

# Solving Code-Related Tasks

Section 2.2

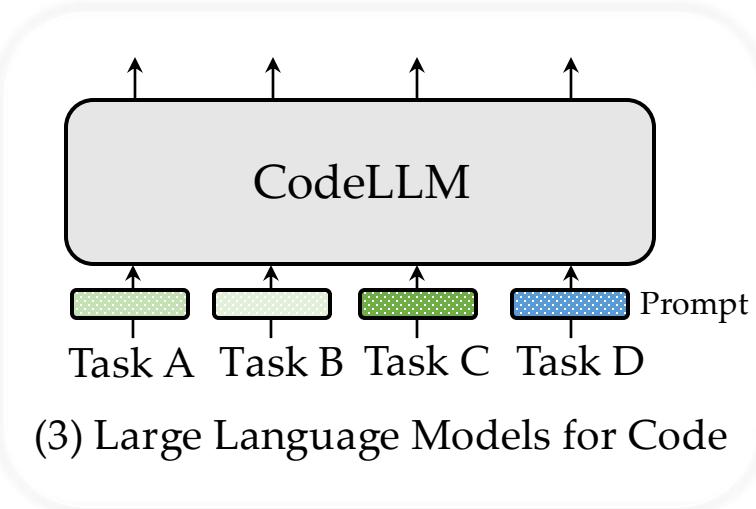
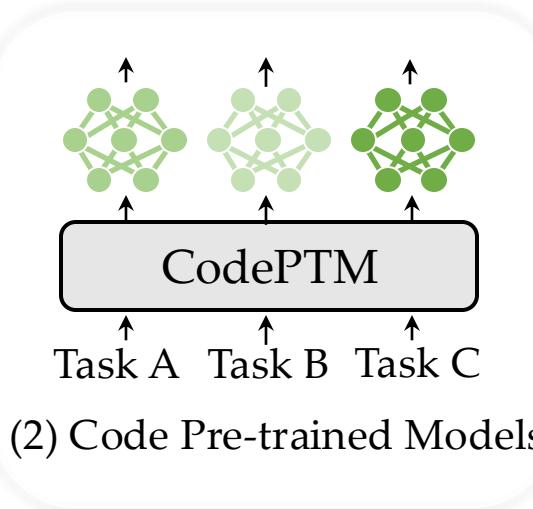
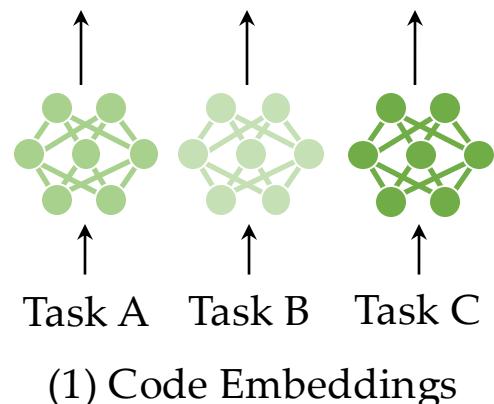
And more tasks ...

Task	Dataset	Date	# PLs.	Description
Clone Detection	POJ-104 [16] <a href="#">[link]</a>	2014	2	a program classification dataset of 52K C/C++ programs
	BigCloneBench [108] <a href="#">[link]</a>	2015	1	a clone detection dataset of eight million Java validated clones
	CLCDSA [109] <a href="#">[link]</a>	2019	3	a cross-language clone dataset of more than 78K solutions
Defect Detection	Devign [78] <a href="#">[link]</a>	2019	1	a dataset of vulnerable C functions
	CrossVul [110] <a href="#">[link]</a>	2021	>40	a dataset of 13K/27K (vulnerable/non-vulnerable) files
	DiverseVul [111] <a href="#">[link]</a>	2023	2	a dataset of 18K/330K (vulnerable/non-vulnerable) functions
Code Repair	Defects4J <a href="#">[link]</a>	2014	1	a database of real Java bugs
	DeepFix [83] <a href="#">[link]</a>	2017	1	a dataset of 7K erroneous C programs for 93 programming tasks
	QuixBugs [112] <a href="#">[link]</a>	2017	2	a multilingual benchmark of similar buggy programs
Code Search	CodeSearchNet [113] <a href="#">[link]</a>	2019	6	a dataset of 6M functions and natural language queries
	AdvTest [114] <a href="#">[link]</a>	2021	1	a Python code search dataset filtered from CodeSearchNet
	WebQueryTest [114] <a href="#">[link]</a>	2021	1	a testing set of Python code search of 1K query-code pairs
Code Translation	CodeTrans [114] <a href="#">[link]</a>	2021	2	a C#/Java dataset collected from several repos
	CoST [115] <a href="#">[link]</a>	2022	7	a dataset containing parallel data from 7 programming languages
	CodeTransOcean [116] <a href="#">[link]</a>	2023	45	a large-scale comprehensive benchmark for code translation
Code Completion	GitHub Java Corpus [2] <a href="#">[link]</a>	2013	1	a giga-token corpus of Java code from a wide variety of domains
	Py150 [117] <a href="#">[link]</a>	2016	1	a corpus of Python programs from GitHub
	LCC [118] <a href="#">[link]</a>	2023	3	a benchmark of code completion with long code context
Code Summarization	CODE-NN [119] <a href="#">[link]</a>	2016	2	a dataset of (title, query) pairs from StackOverflow
	TL-CodeSum [120] <a href="#">[link]</a>	2018	1	a dataset containing 69K pairs of (API sequence, code, summary)
	CodeSearchNet [113] <a href="#">[link]</a>	2019	6	a dataset of 6M functions and natural language queries
GitHub	CommitGen [121] <a href="#">[link]</a>	2017	4	a multilingual dataset collected from open source projects
	CommitBERT [122] <a href="#">[link]</a>	2021	6	a multilingual dataset of code modification and commit messages
	SWE-bench [123] <a href="#">[link]</a>	2023	1	a benchmark of 2K SE problems and corresponding PRs

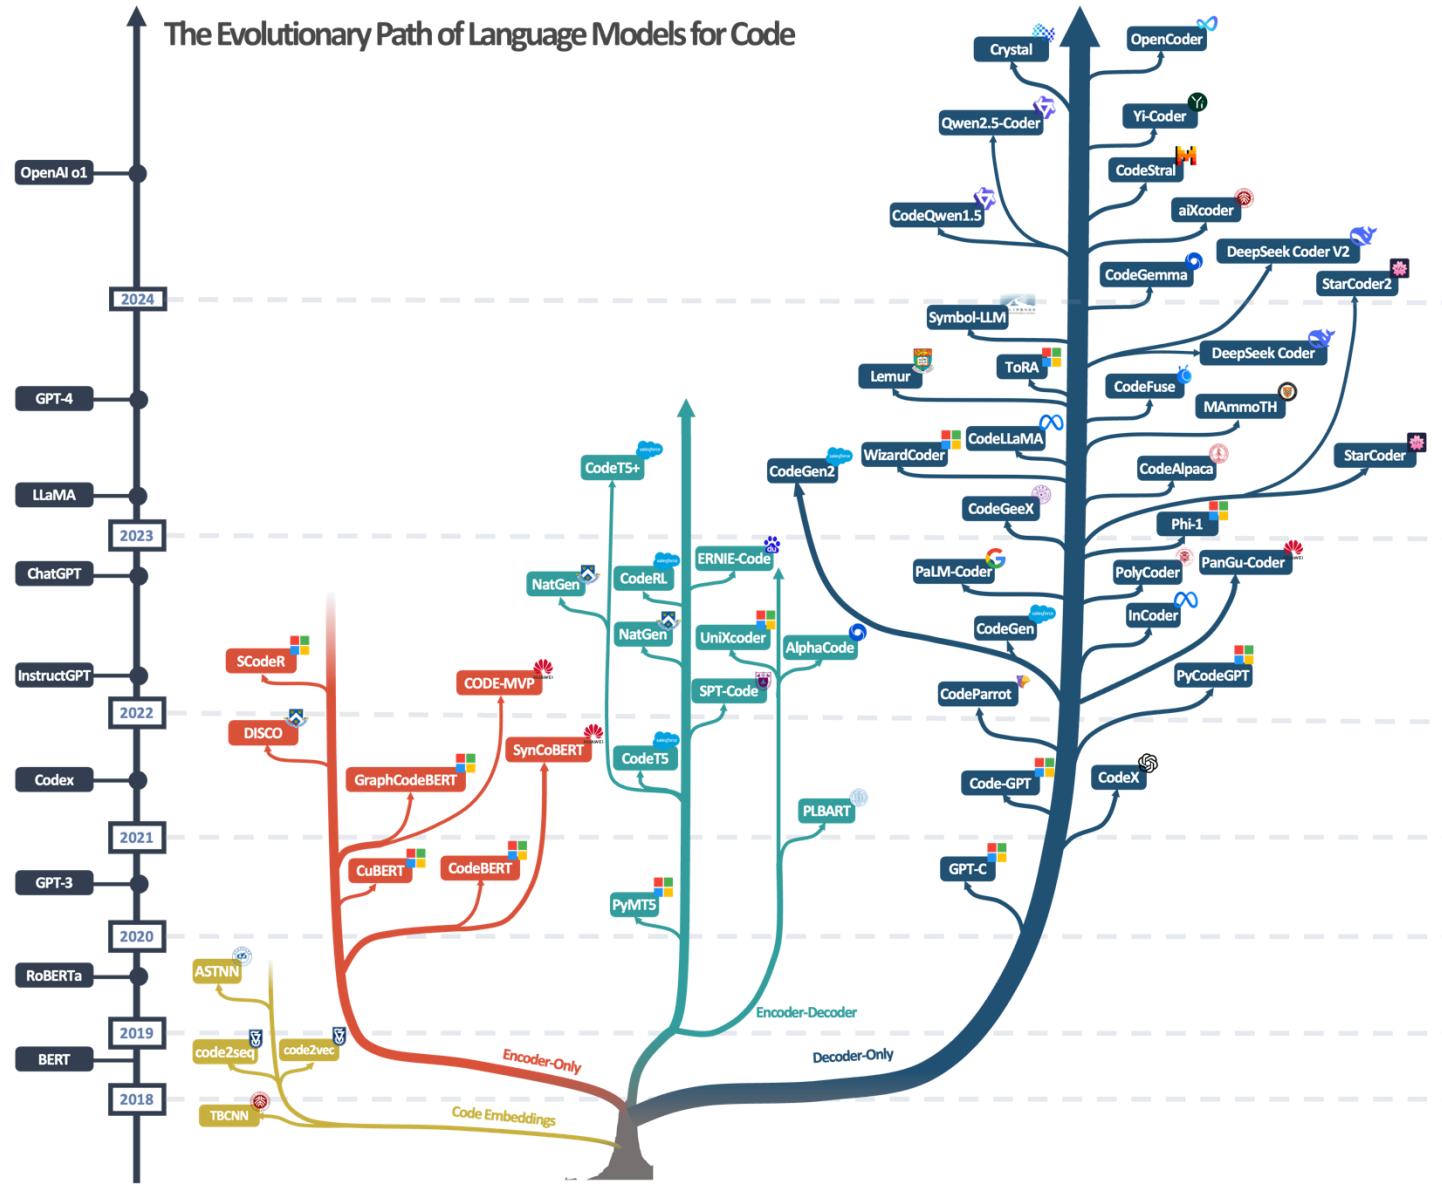
Well, essentially, they all require **task-specific modeling from scratch**.

# The Paradigm Shift of CodeLMs

Section 4



# The evolution from the perspective of models



# Code Pre-trained Models (CodePTMs)

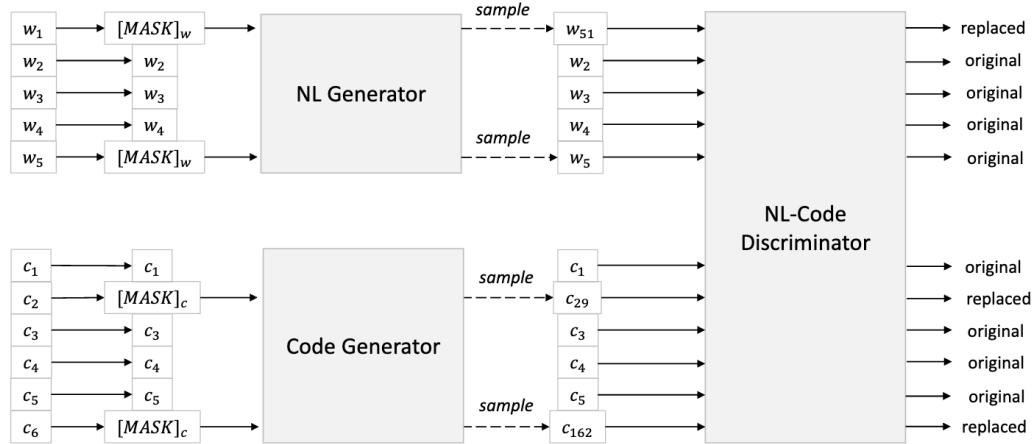
Section 3.1



# CuBERT and CodeBERT

Section 3.1

## CodeBERT: A Pre-Trained Model for Programming and Natural Languages



CodeBERT: A code version of RoBERTa

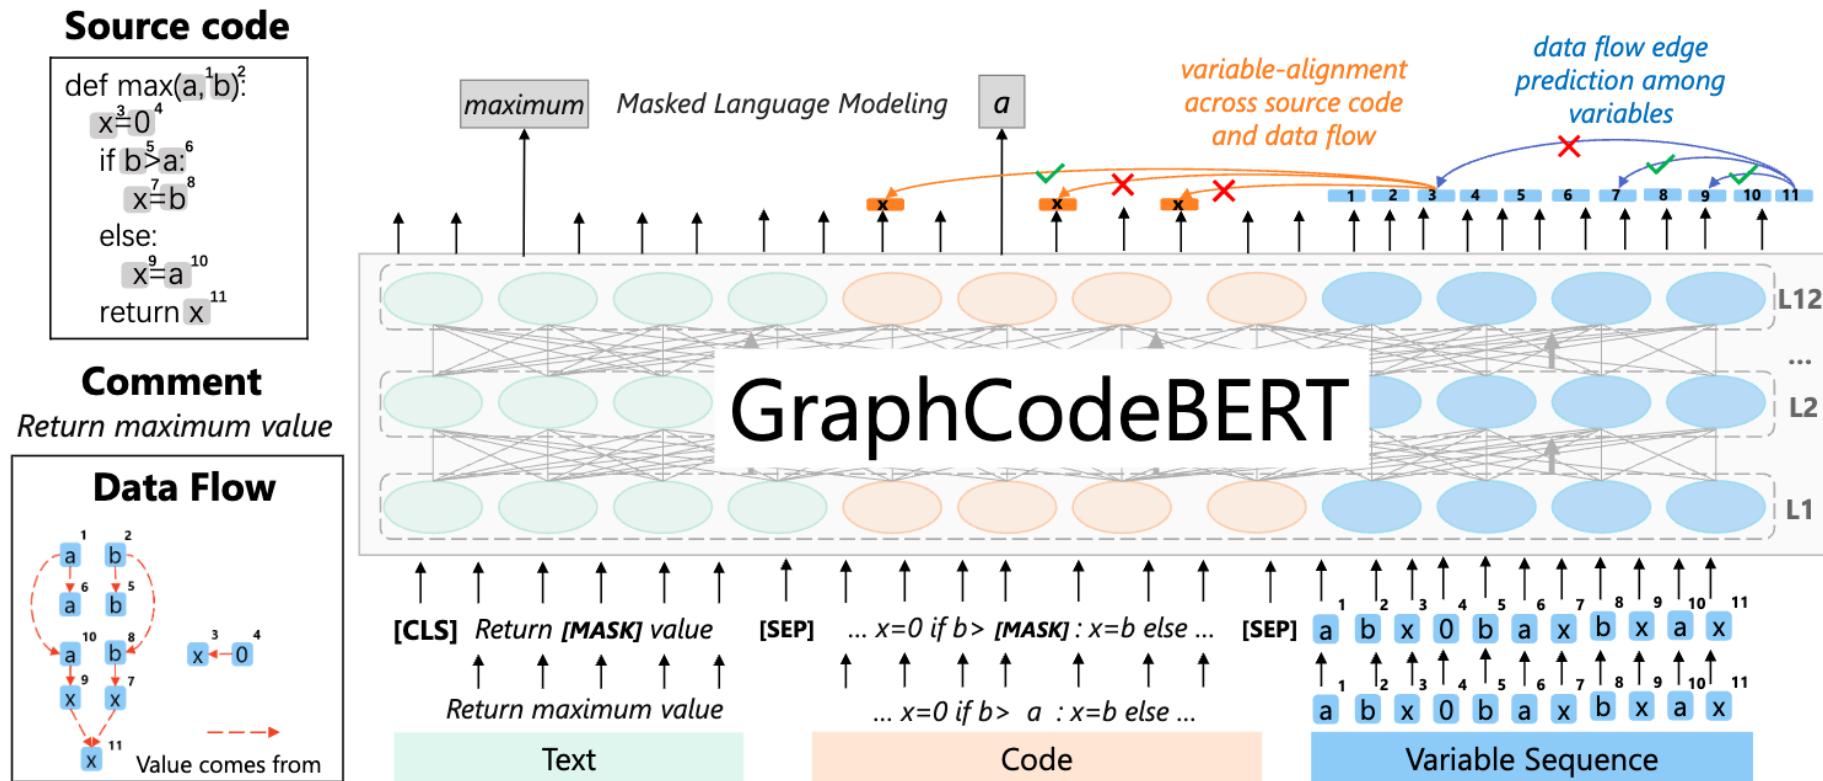
## Learning and Evaluating Contextual Embedding of Source Code

CuBERT: A code version of BERT

# GraphCodeBERT

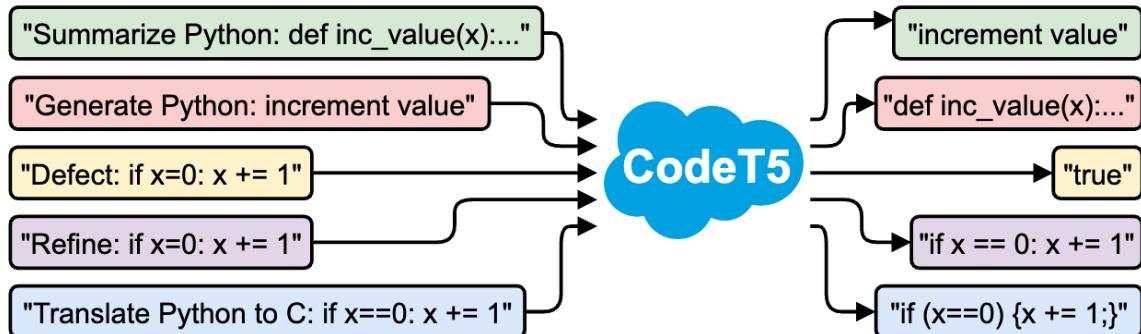
Section 3.1

How about typical “BERT-Style” Training meets Code Structures



# T5 and BART like CodePTMs

Section 3.1



PLBART Encoder Input	PLBART Decoder Output
Is 0 the [MASK] Fibonacci [MASK] ? <En>	<En> Is 0 the <b>first</b> Fibonacci <b>number</b> ?
public static main ( String args [ ] ) { date = Date () ; System . out . ( String . format ( " Current Date : % tc " , )) ; } <java>	<java> public static <b>void</b> main ( String args [ ] ) { <b>Date</b> date = <b>new</b> Date () ; System . out . <b>printf</b> ( String . format ( " Current Date : % tc " , <b>date</b> )) ; }
def addThreeNumbers ( x , y , z ) : NEW_LINE INDENT return [MASK] <python>	<python> def addThreeNumbers ( x , y , z ) : NEW_LINE INDENT return x + y + z

## PLBART: Denoising Pre-training

1. Standard denoising training for T5 and BART models
2. Identifier types can be used for sequence labeling learning
3. Seq2seq learning unique to code
  - Deobfuscation
  - Naturalization
  - Mutual generation of code and comments

CodeT5: Identifier-aware Unified Pre-trained Encoder-Decoder Models for Code Understanding and Generation, EMNLP 2021

Unified Pre-training for Program Understanding and Generation, NAACL 2021

# Some issues

- When introducing code features, changes to the vocabulary, input format, or attention patterns often prevent **generalization**.
- The **generation capability** is quite weak.
- When adapting to downstream tasks, **fine-tuning** is typically required.

Architecture	Models	Struct.	Base	Strategy	Size
Encoder	CuBERT [179]	✗	-	MLM + NSP	340M
	CodeBERT [28]	✗	RoBERTa	MLM + RTD	125M
	GraphCodeBERT [182]	✓	CodeBERT	MLM + Edge Pred. + Node Align.	125M
	SynCoBERT [184]	✓	CodeBERT	MMLM + IP + TEP + MCL	125M
	CODE-MVP [185]	✓	GraphCodeBERT	FGTI + MCL + MMLM	125M
	SCoDeR [187]	✓	UniXcoder	Soft-Labeled Contrastive Pre-training	125M
	DISCO [186]	✓	-	MLM + NT-MLM + CLR	110M
Enc-Dec	PLBART [30]	✗	-	Denoising Pre-training	140M/406M
	CodeT5 [29]	✓	-	MSP + IP + MIP + Bimodal Generation	60M/220M/770M
	PyMT5 [194]	✗	-	MSP	374M
	UniXcoder [195]	✓	-	MLM + ULM + MSP + MCL + CMG	125M
	NatGen [196]	✓	CodeT5	Code Naturalization	220M
	TreeBERT [192]	✓	-	TMLM + NOP	210M
	ERNIE-Code [197]	✗	mT5	SCLM + PTLM	560M
	CodeExecutor [198]	✗	UniXcoder	Code execution + Curriculum Learning	125M
	LongCoder [118]	✗	UniXcoder	CLM	150M
Decoder	GPT-C [199]	✗	-	CLM	366M
	CodeGPT [114]	✗	-	CLM	124M
	PyCodeGPT [200]	✗	GPT-Neo	CLM	110M

# If you want to learn more ...

Check out this post I wrote during my UG thesis in 2022!

知乎 | 首发于 Paper Reading

... 写文章



## [代码表征] Code预训练语言模型学习指南（原理/分析/代码）

NaturalSelection ✅  
香港大学 计算机科学博士在读

雨打蕉叶潇潇几夜、sonta 等 113 人赞同了该回答 >

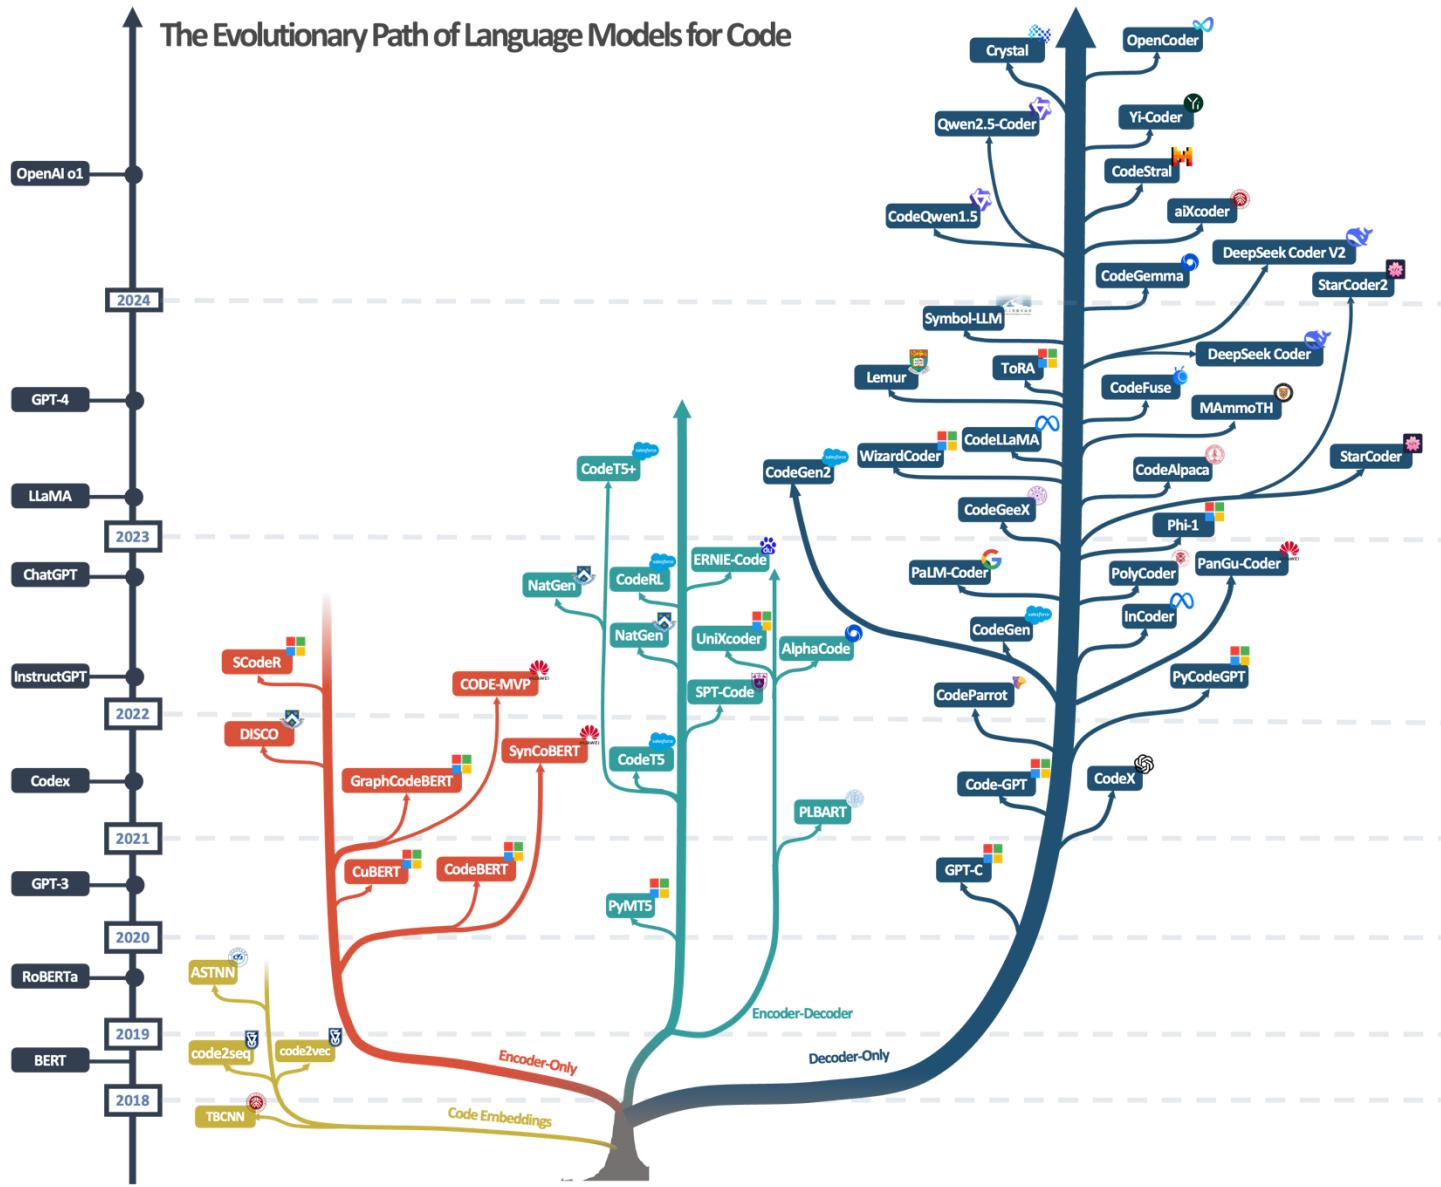
自从2020年CodeBERT开了代码表征预训练模型（本文称之为CodePTM）这个新坑后，在短短两年的时间内出现了若干个程序设计语言（Programming Language，称之为PL，与Natural Language，也就是NL对应）语言模型。它们的共同特点是大部分用于处理PL（或Source Code）

▲ 赞同 113 ● 6 条评论 ❤ 喜欢 ★ 收藏 ○ 设置

<https://zhuanlan.zhihu.com/p/539929943>

# Large Language Models for Code

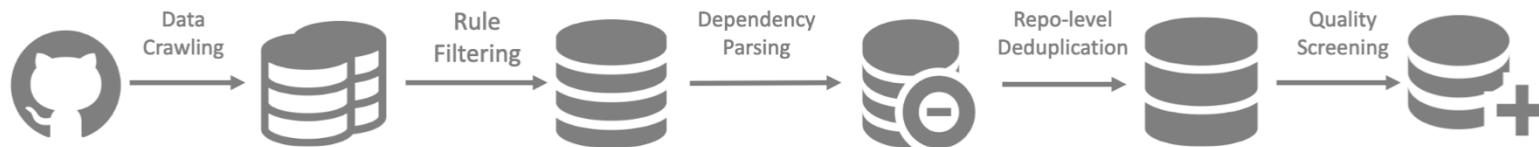
Section 4.1



# Large Language Models for Code

Section 4.1

At early stage, training from scratch.



DeepSeek Coder V1

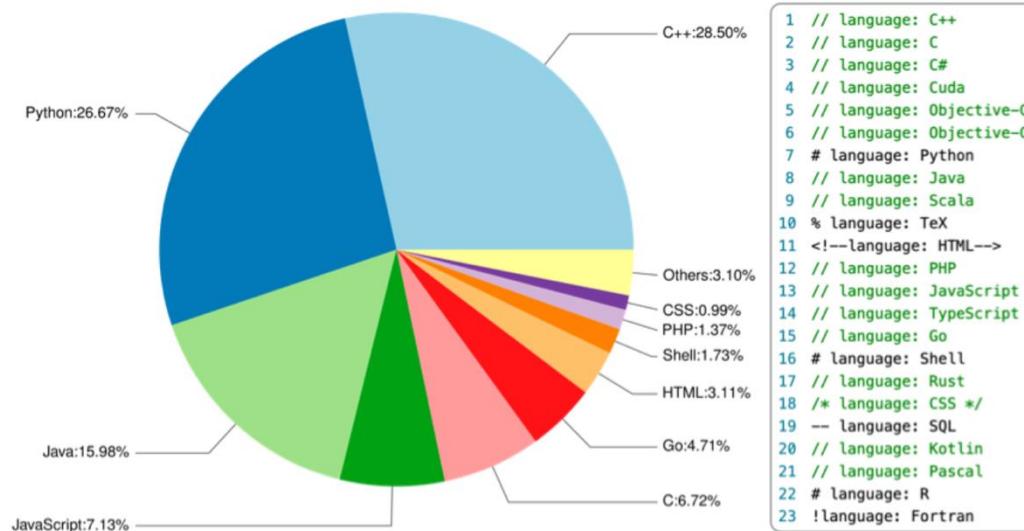


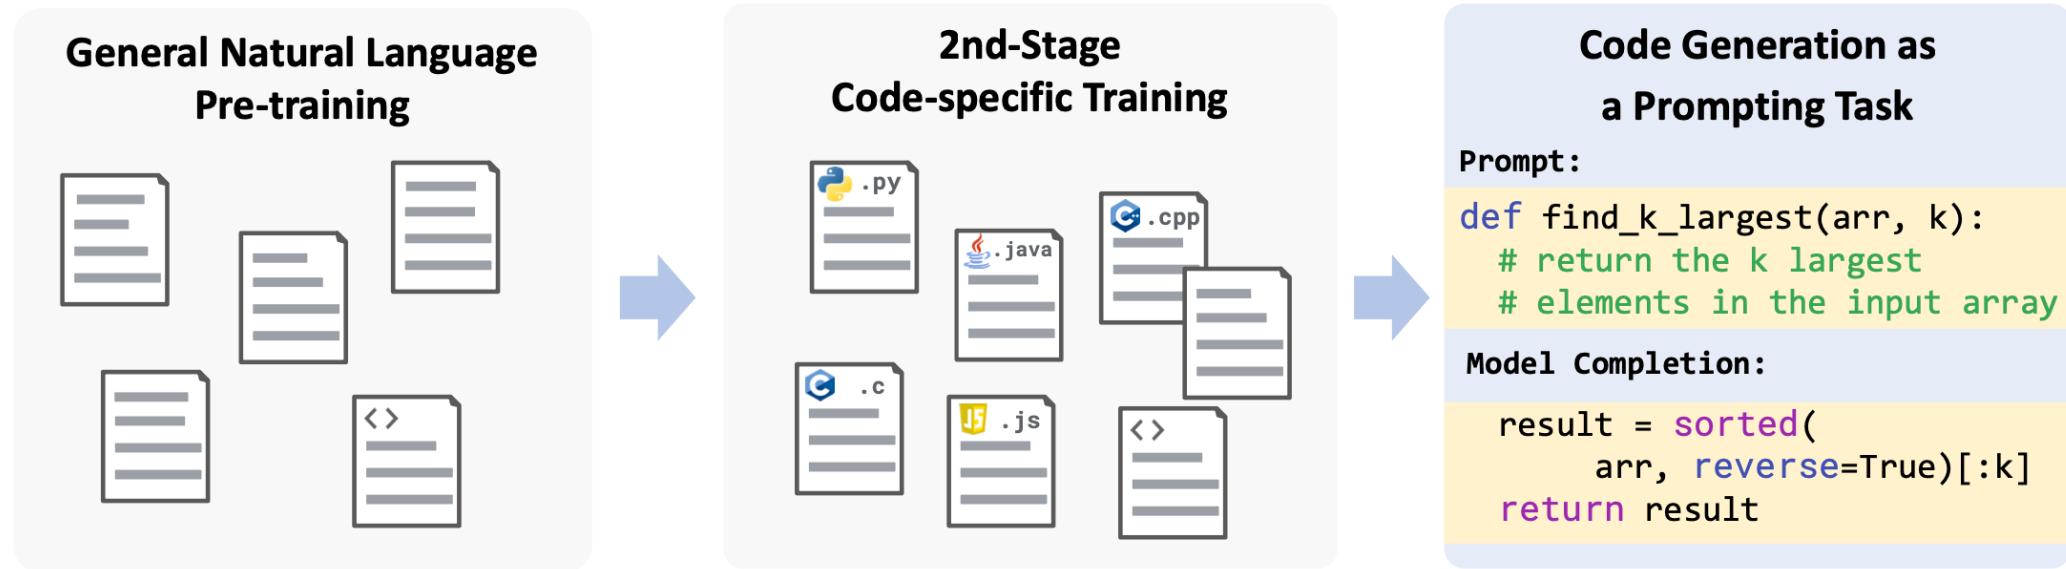
Figure 3: Language distribution and tags of CodeGeeX's data.

CodeGeeX V1

# Large Language Models for Code

Section 4.1

Gradually...



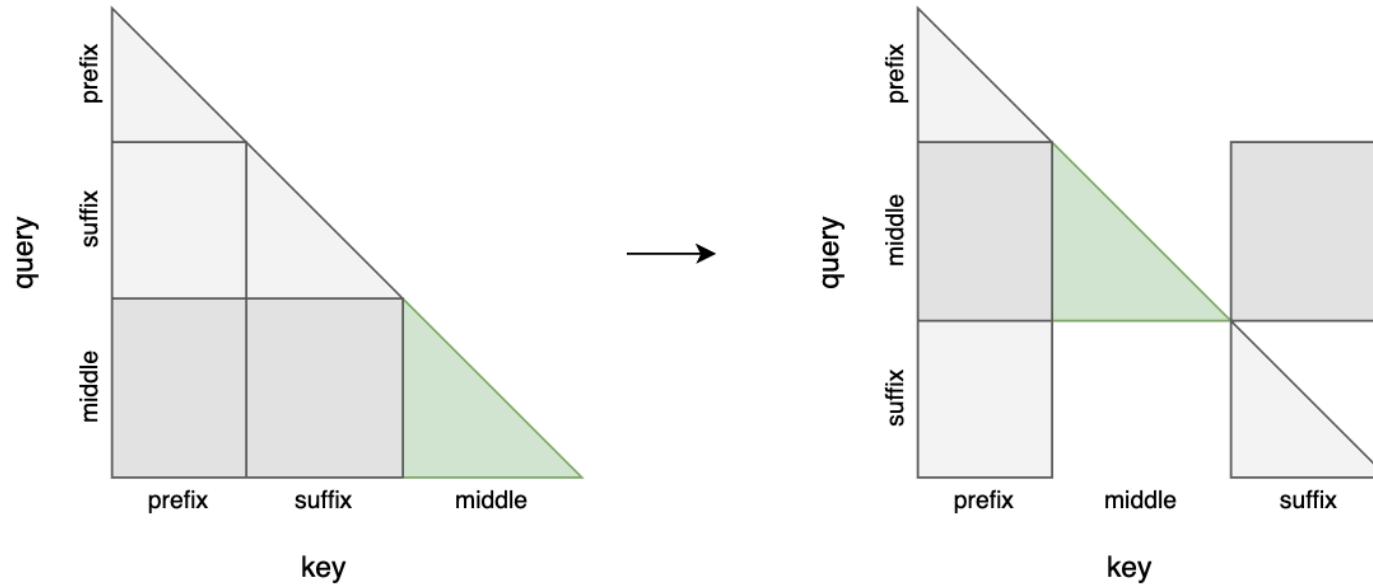
Other more-recent Code LLMs:

- Code LLaMA 
- DeepSeek Coder 

# FIM Training for CodeLLMs

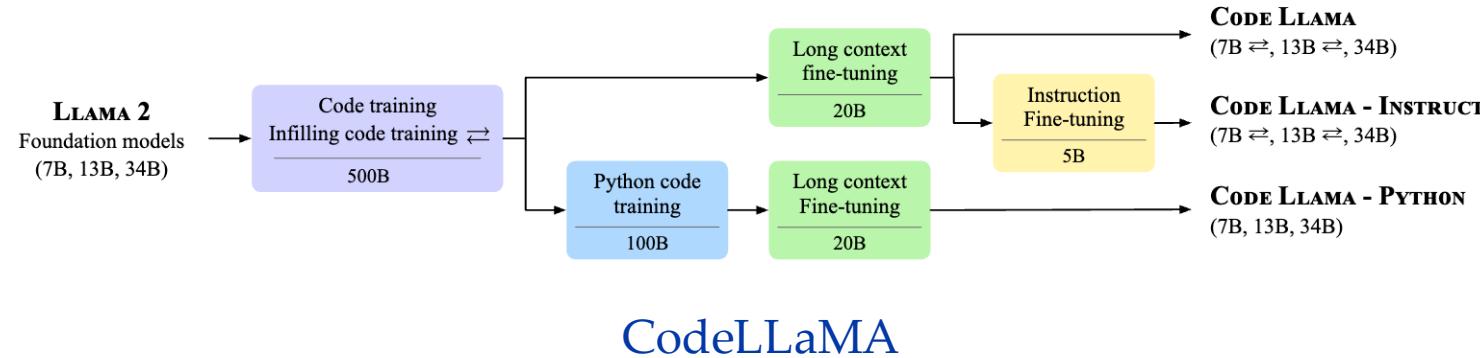
Section 4.1

- Consistent with AR training, like GPT
- FIM (Fill-in-the-Middle) pretraining

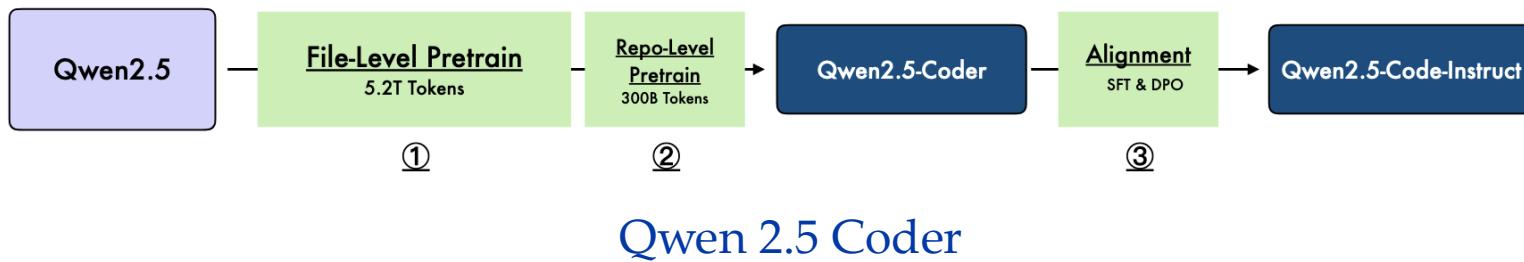


# Large Language Models for Code

Section 4.1



CodeLLaMA



Qwen 2.5 Coder

And Codex, PaLM Coder, DeepSeekCoder V2 ...

# Large Language Models for Code

Section 4

Arch.	Model Name	Size	Base	Vocab.	Context	Training Objs.	Data Scale	Public
Enc-Dec	AlphaCode [275]	284M/1.1B/ 2.8B/8.7B/ 41.1B	-	8.0K	1536+768	MLM+CLM	354B/590B/ 826B/1250B/ 967B	✗
	CodeT5+ [276]	220M/770M/ 2B/6B/16B	CodeGen	50.0K	2048+2048	MSP+CLM+CL	51.5B	✓
Decoder / CodeLLMs	Codex [36]	2.5B/12B	-	50.3K	4K	CLM	100B/159GB	✓
	CodeParrot [288]	125M/1.5B	-	32.8K	1K	CLM	26B/50GB	✓
	PolyCoder [289]	160M/0.4B/2.7B	-	50.3K	2K	CLM	39B/254GB	✓
	CodeGen [277]	350M/2.7B/ 6.1B/16.1B	-	50.0K	2K	CLM	1.2T	✓
	PaLM-Coder [34]	8B/62B/540B	PaLM	256K	2K	CLM	7.75B	✗
	InCoder [94]	1.3B/6.7B	-	50.3K	2K	FIM	52B/159GB	✓
	PanGu-Coder [290]	317M/2.6B	-	42K	1K	CLM+MLM	387B/147GB	✗
	SantaCoder [291]	1.1B	-	49.2K	2K	FIM	236B/268GB	✓
	phi-1 [292]	350M/1.3B	-	50.0K	2K	CLM	7B	✓
	CodeGeeX [293]	13B	-	52.2K	2K	CLM	850B	✓
	CodeGen2 [286]	1B/3.7B/7B/16B	-	50.0K	2K	MLM+CLM	400B	✓
	StarCoder [279]	15.5B	-	49.2K	8K	FIM	1T/815GB	✓
	CodeAlpaca [294]	7B/13B	LLaMA	32.0K	4K	CLM	20K	✓
	WizardCoder [295]	1B/3B/7B/ 13B/15B/34B	StarCoder	32.0K	2K	CLM	78k	✓
	AquilaCode [296]	7B	Aquila	100.0K	2K	CLM	-	✓
	CodeGeeX2 [293]	6B	ChatGLM2	65.0K	8K	CLM	600B	✓
	CodeLLaMA [297]	7B/13B/34B/70B	LLaMA2	32.0K	4K	FIM	500B	✓
	ToRA-Code [298]	7B/13B/34B	CodeLLaMA	32.0K	2K	Min. NLL	223K	✓
	MAmmoTH-Coder [299]	7B/13B/34B	CodeLLaMA	32.0K	2K	CLM	260K	✓
	Code-Qwen [300]	7B/14B	Qwen	152.0K	8K	CLM	90B	✓
	CodeFuse [301]	1.3B/6.5B/ 13B/34B	Multiple	100.9K	4K	CLM	1.6TB	✓
	CodeShell [302]	7B	-	70.1K	8K	CLM	500B	✓
	Lemur [303]	70B	LLaMA2	32.0K	4K	CLM	90B	✓
	DeepSeekCoder [304]	1.3B/5.7B/ 6.7B/33B	-	32.0K	16K	FIM	2T	✓
	Symbol-LLM [305]	7B/13B	LLaMA2	32.0K	4K	CLM	2.25GB	✓
	Stable Code [306]	3B	-	50.3K	16K	FIM	1.3T	✓
	DeciCoder [307]	1B/6B	-	49.2K	2K	FIM	446B	✓
	StarCoder2 [280]	3B/7B/15B	-	49.2K	16K	FIM	900B/3TB	✓
	CodeGemma [308]	2B/7B	Gemma	250K	8K	FIM	1T	✓
	CodeStral [309]	22B	-	32.0K	32k	CLM+FIM	-	✓
	DeepSeekCoderV2 [310]	16B/236B	DeepSeekV2	100K	128K	CLM+FIM	10.2T	✓
	Crystal [311]	7B	-	32K	2K	CLM	1.4T	✓
	Yi-Coder [312]	1.5B/9B	Yi	64K	128K	CLM	2.4T	✓
	OpenCoder [313]	1.5B/8B	-	96.6K	8K	CLM	2.5T	✓

# General-purpose Code Generation

Section 4

HumanEval Doc-string2Code (Chen et al., 2021)

```
def sum_odd_elements(lst):
    """given non-empty list of integers, return the
    sum of all of the odd elements that are in even
    positions

Examples
solution([5, 8, 7, 1]) => 12
solution([3, 3, 3, 3, 3]) => 9
solution([30, 13, 24, 321]) => 0
"""
return sum([
    lst[i] for i in range(0, len(lst))
    if i % 2 == 0 and list[i] % 2 == 1])
```

MBPP NL description + tests (Austin et al., 2021)

```
Write a function to find the smallest missing
element in a sorted array. Your code should
satisfy these tests:

assert smallest_missing(
    [0, 1, 2, 3, 4, 5, 6], 0, 6) == 7
assert smallest_missing(
    [0, 1, 2, 6, 9, 11, 15], 0, 6) == 3
assert smallest_missing(
    [1, 2, 3, 4, 6, 9, 11, 15], 0, 7) == 0

def smallest_missing(arr, n, m):
    smallest = min(n, m)
    for i in range(n, m + 1):
        if arr[i] <= smallest:
            smallest += 1
    return smallest
```

Python Algorithmic Problems

# Competition Level Programming

Section 4

Problem	Generated Code	Test Cases
<p>H-Index</p> <p>Given a list of citations counts, where each citation is a nonnegative integer, write a function <code>h_index</code> that outputs the h-index. The h-index is the largest number <math>h</math> such that <math>h</math> papers have at least <math>h</math> citations.</p> <p>Example:</p> <p>Input: [3,0,6,1,4] Output: 3</p>	<pre>def h_index(counts):     n = len(counts)     if n &gt; 0:         counts.sort()         counts.reverse()         h = 0         while (h &lt; n and                counts[h]-1&gt;=h):             h += 1         return h     else:         return 0</pre>	<p>Input: [1,4,1,4,2,1,3,5,6]</p> <p>Generated Code Output: 4 ✓</p> <p>Input: [1000,500,500,250,100, 100,100,100,100,75,50, 30,20,15,15,10,5,2,1]</p> <p>Generated Code Output: 15 ✓</p>

An example competition-level coding problem (figure from Hendrycks et al. 2021)

APPS and CodeContests

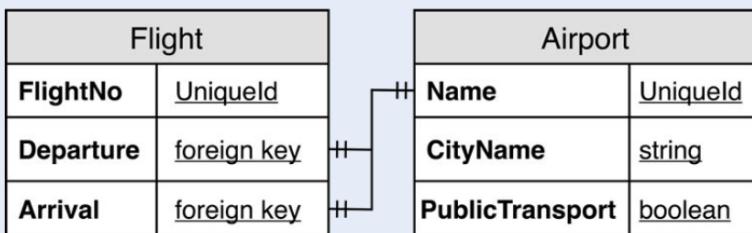
# Code Generation to Domain-Specific Programs

Section 4

## Natural Language Questions with Database Schema

### Input Utterance

Show me flights from Pittsburgh to SFO



### SQL Query

```
SELECT Flight.FlightNo
FROM Flight
JOIN Airport as DepAirport
ON
    Flight.Departure == DepAirport.Name
JOIN Airport as ArvAirport
ON
    Flight.Arrival == ArvAirport.Name
WHERE
    DepAirport.CityName == Pittsburgh
    AND
    ArvAirport.CityName == San_Francisco
```

Text-to-SQL

# Large Language Models for Code

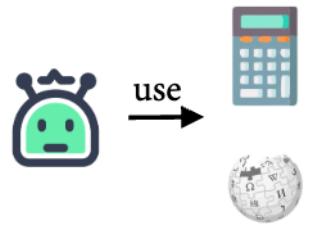
Section 4

The coding ability has obviously become stronger

but **pure code training** often sacrifices other performance aspects,  
making the model impractical in the real world. For example:

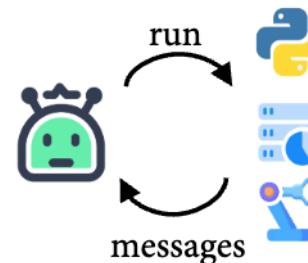
## Augment with Tools

Employ various tools  
to augment agents' capabilities



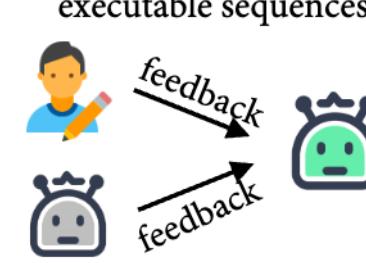
## Self-Debug

Utilize error messages from the environment to rectify existing errors



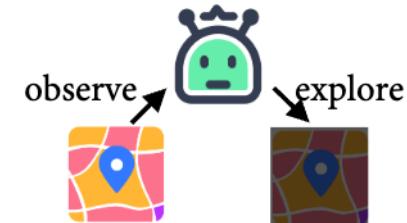
## Follow Feedback

Understand complex feedback from human / agents and convert them into symbolic executable sequences



## Explore Environment

Operate in environments that are partially observable



# Balancing Coding and NL

Section 4.1

## For Agentic Use



LLaMA2

Code-Centric  
Pre-training  
→  
90B Tokens



Lemur

SFT  
→  
300K Tokens



Lemur-Chat

## For Efficiency

NL

Phase1: SlimPajama



CRYSTAL 7B

Phase2: 63% code from Stack

Phase2: 37% NL

NL + Code

# CodeLLMs as Base

Section 5.2

Early



CodeGen-16B-mono

Chinese Tokens Pre-training



MOSS

Recent



DeepSeek-Coder-Base-v1.5 7B

120B Math Tokens

NL and Code Data



DeepSeekMath 7B

MOSS: An Open Conversational Large Language Model, Machine Intelligence Research  
DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models

# CodeLLMs as Base

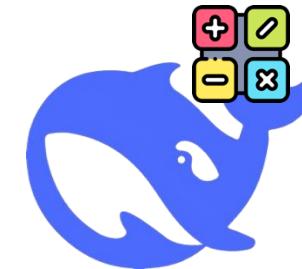
## Some findings

1. Code training or pre-training and tasks like math are largely not mutually exclusive; in fact, they often enhance each other.
2. How to balance their **proportions** is very important.



120B Math Tokens

NL and Code Data



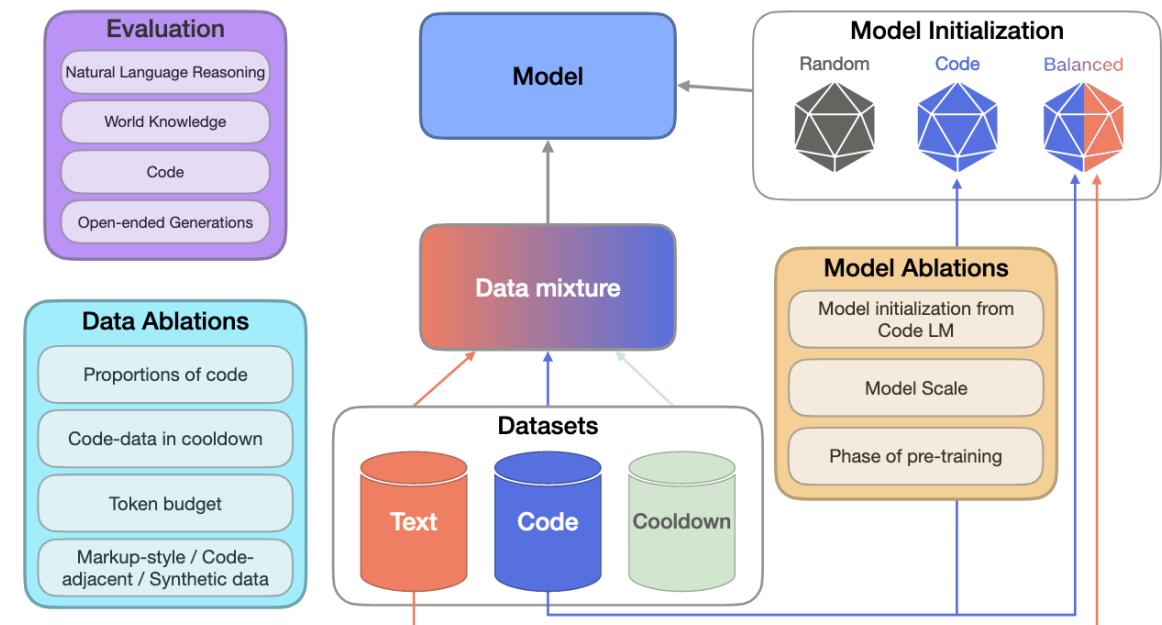
DeepSeek-Coder-Base-v1.5 7B

DeepSeekMath 7B

# CodeLLMs as Base

More findings, beyond math

1. Non-code tasks, performance peaks on average when the code proportion is 25%.
2. Excessive code reduces world knowledge
3. Code performance improves linearly as the code proportion increases.

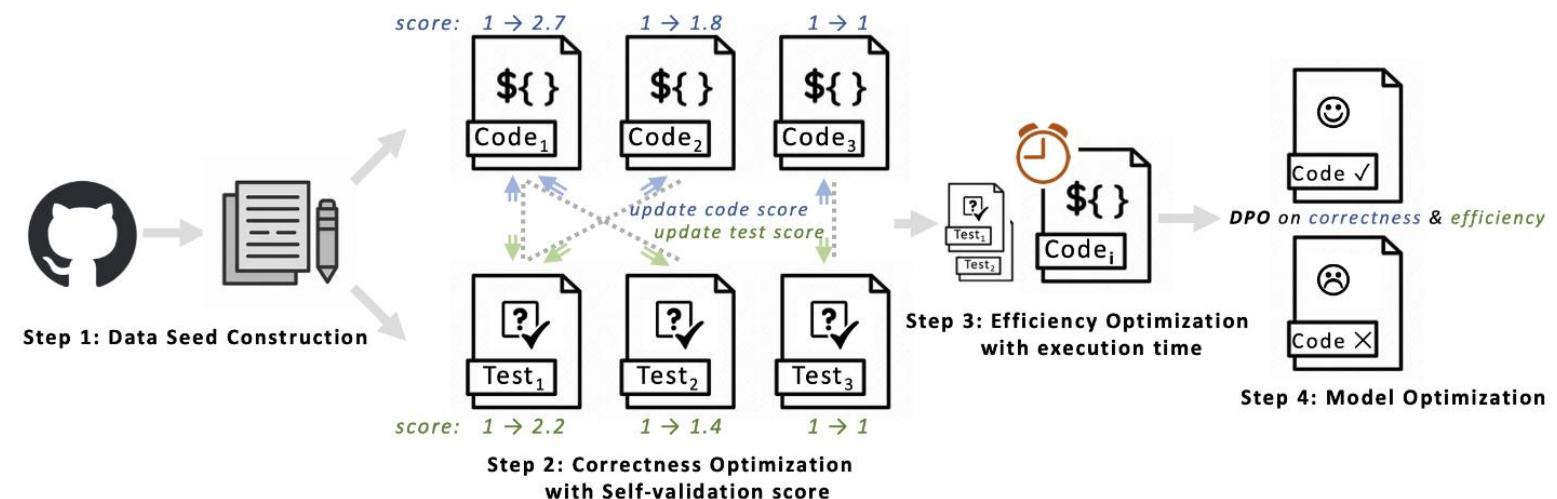
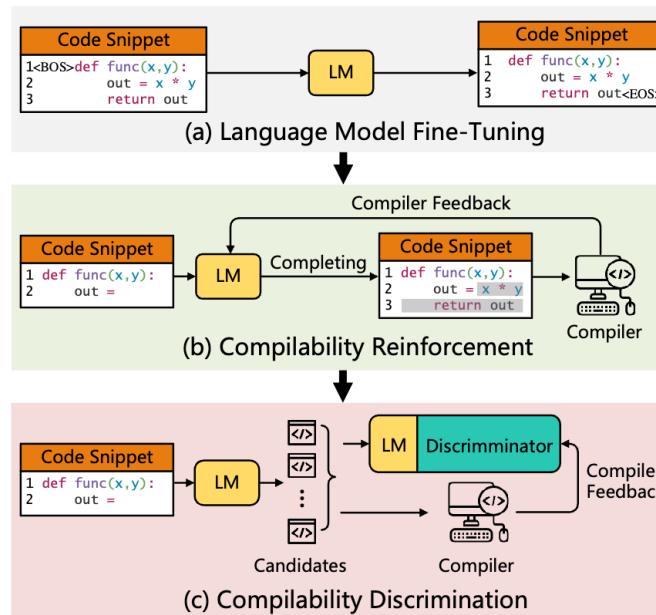


# Preference Optimization + Compiler feedback

Section 4.4

**Insufficient Priority on Correctness:** In ambiguous cases, CodeLLMs fail to prioritize the correct solution over an incorrect one.

**Runtime Efficiency:** The generated code, while functionally correct, may have performance issues).



# The applications of Code Intelligence



Software  
Engineering



Reasoning

AI4Science



## The Application of Code Intelligence

Agents

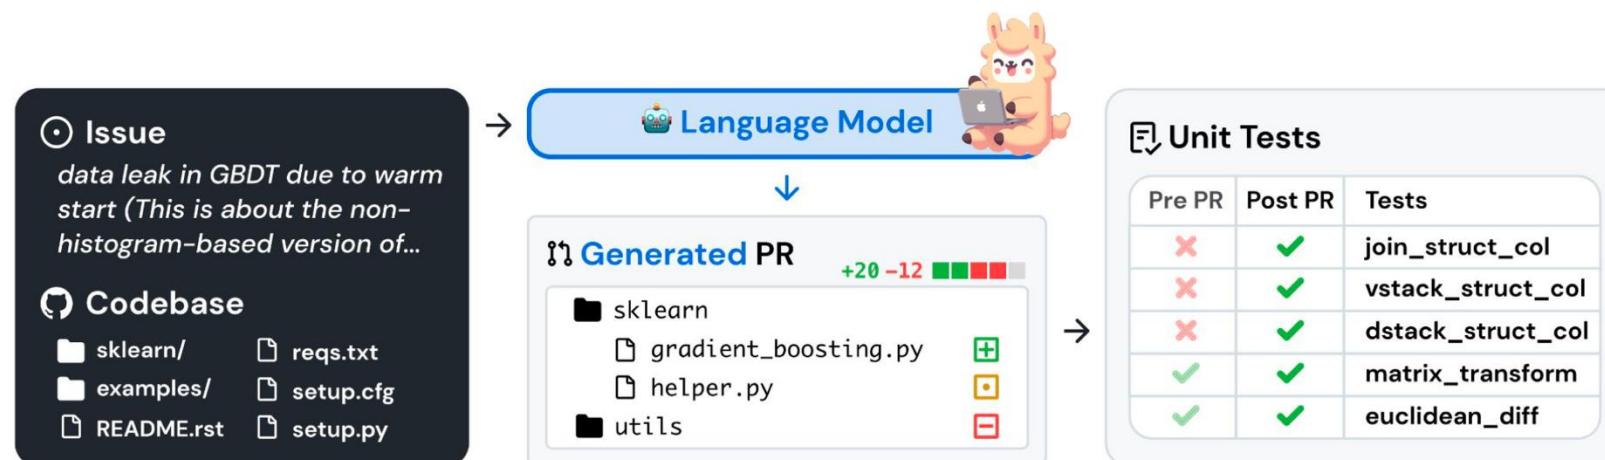


# Applications: Software Engineering

Section 6.1

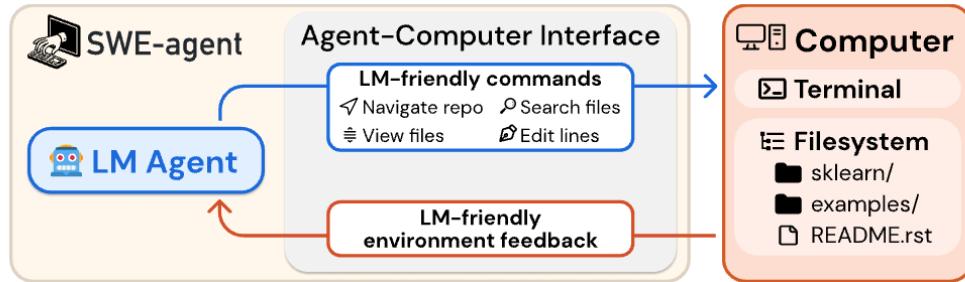
Move beyond simple code generation

Real PRs from popular Python open-source repositories (e.g., Django, Flask, etc.), ultimately filtering out valid task instances. Each task instance corresponds to a GitHub Issue and its merged solution.



# Applications: Software Engineering

Section 6.1



An LM interacting with a computer through an agent-computer interface

SWE-agent lets your LM autonomously use tools to:

1. Fix issues in **real GitHub repositories**,
2. perform tasks on the **web**,
3. find cybersecurity vulnerabilities (by solving Capture The Flag challenges),
4. Custom Tasks

# Applications: Software Engineering

Section 6.1

We are witnessing “Full-Stack” SE platforms

1. Modify code
2. Run commands
3. Browse the web
4. Call APIs
5. Even copy code snippets from StackOverflow.

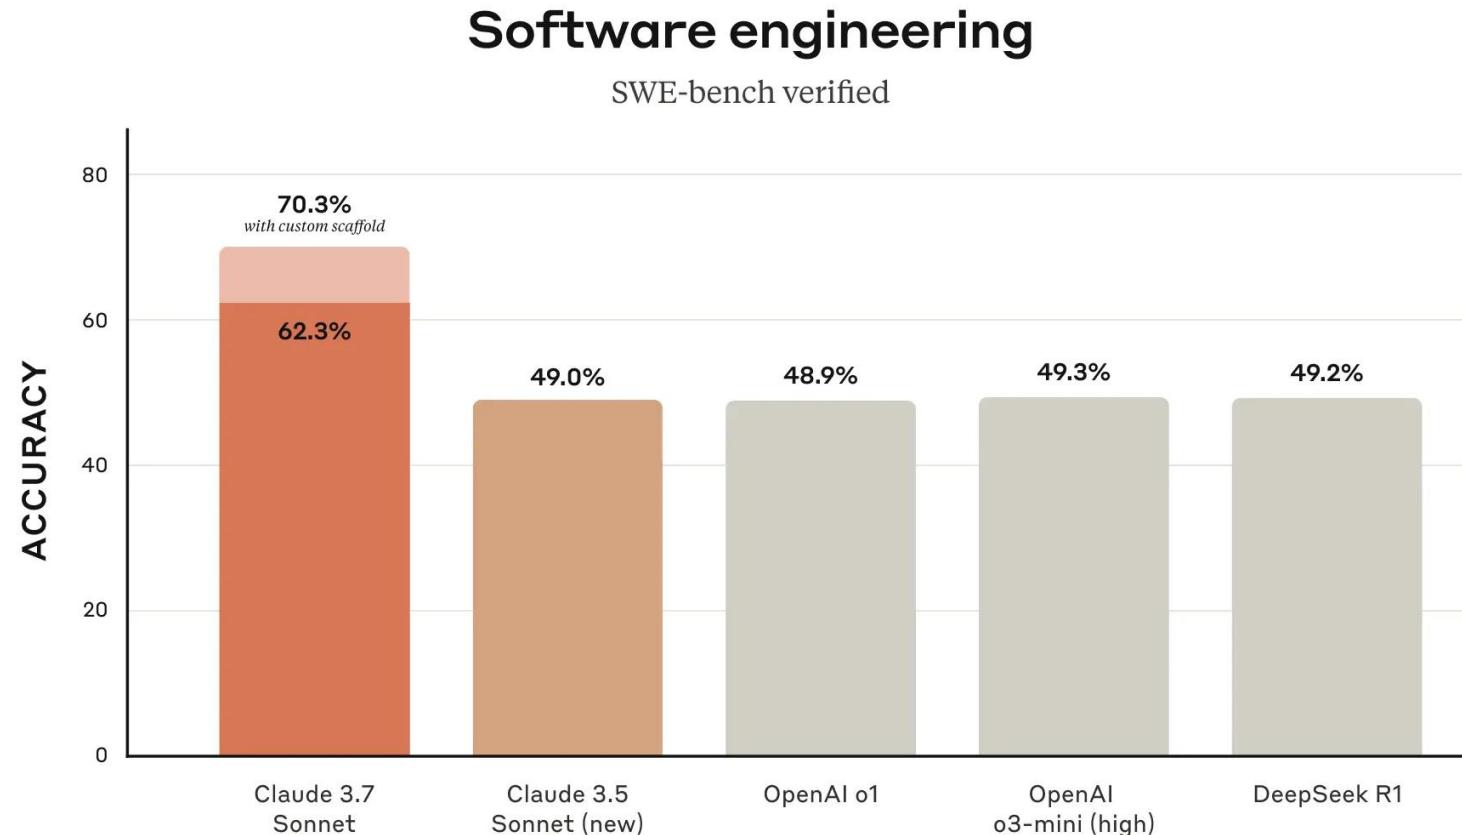


OpenDevin / OpenHands

# Applications: Software Engineering

Section 6.1

More advanced models + frameworks Are “dominating” these benchmarks.



1. Claude 3.7 Sonnet
2. Scaffolding

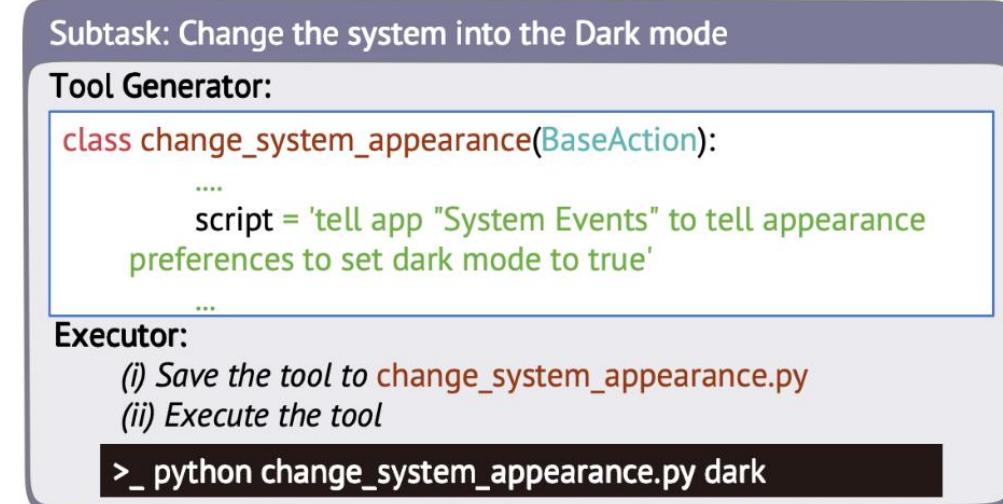
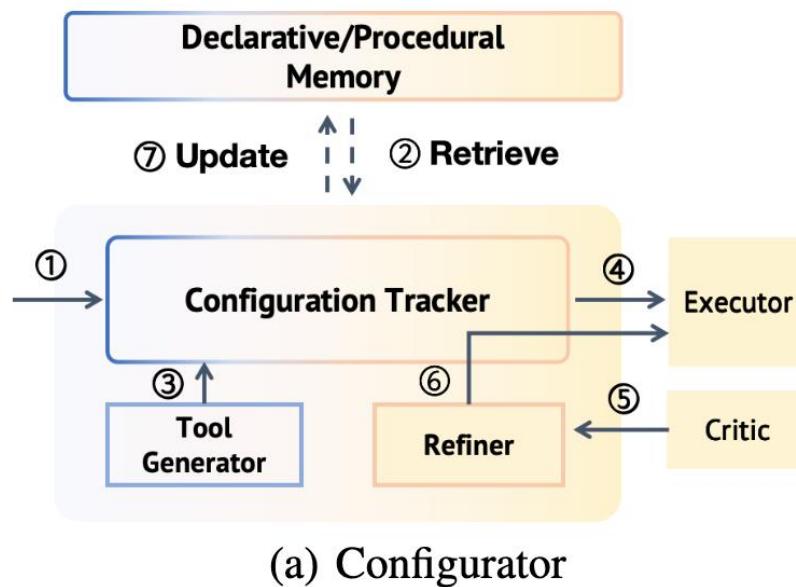
Agentic coding

# Applications: Agents

Section 6.3



OS-Copilot: -> Code-based computer agents Framework



(b) A running example

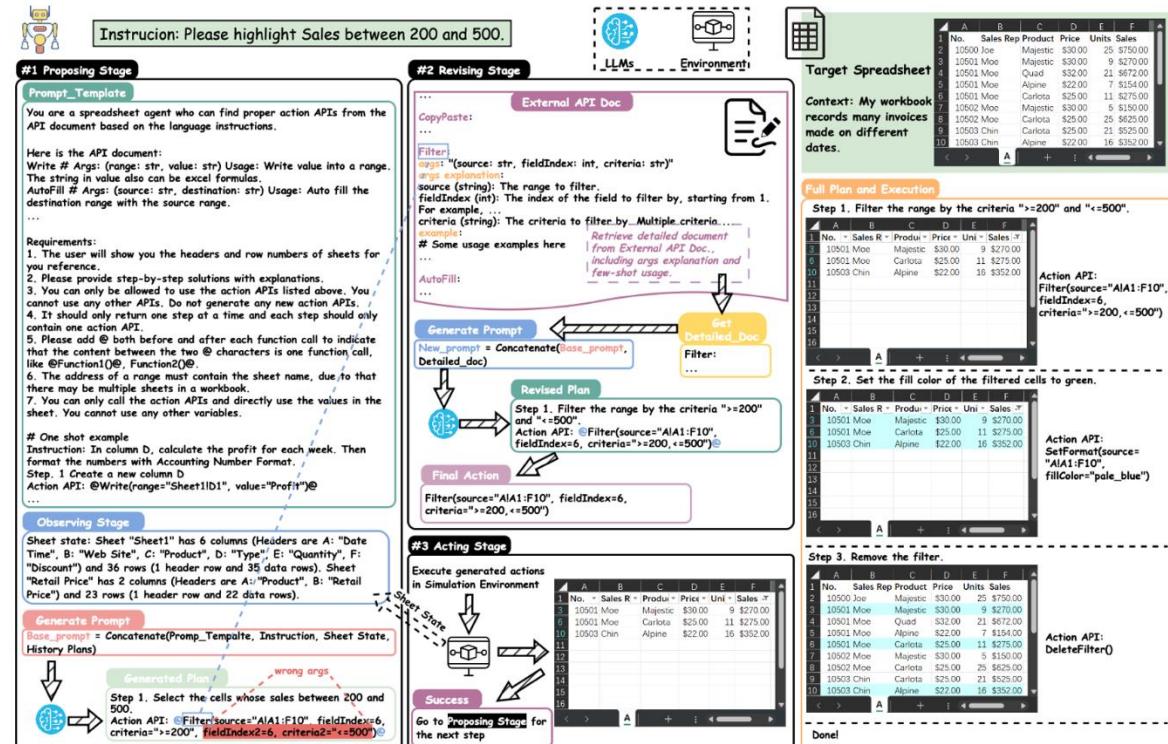
# Applications: Agents

Section 6.3

What kind of issues these agents can solve?

1. API-interface available
2. CLI, like “Apple Script”
3. Numerical Calculataions

Generate code + Invoke API -> Solve computer task



A case in SheetCopilot

# Applications: Agents

Section 6.3

Use executable code to consolidate LLM agents' actions into a unified action space

Integrated with a Python interpreter, execute **code actions** and **dynamically revise prior actions** or **emit new actions** upon new observations (e.g., code execution results) through multi-turn interactions

**Example**

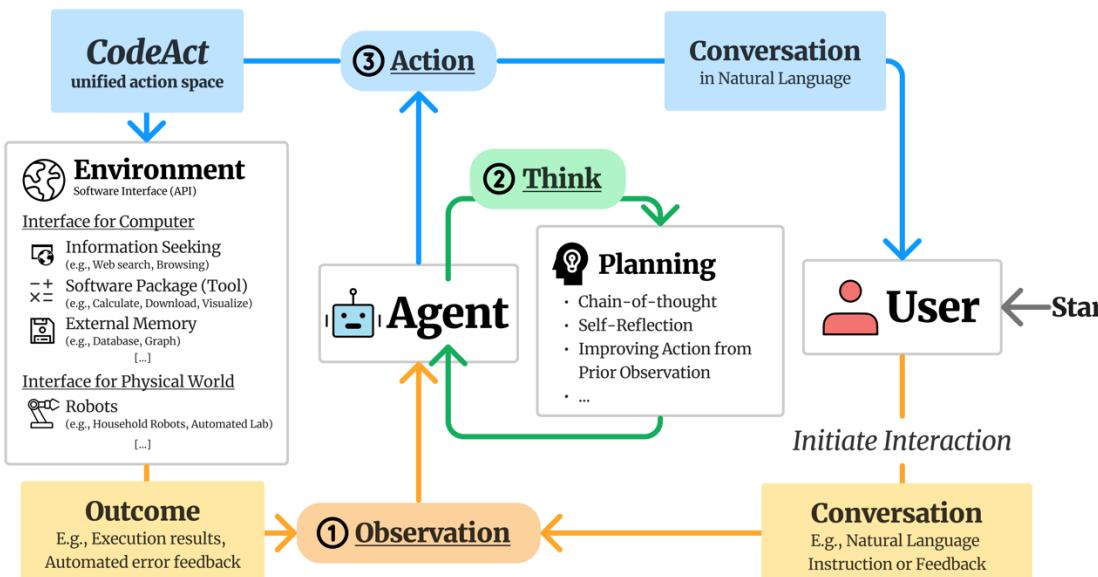
User:  
① Find the sum of the reciprocals of the roots of  $x^2 - 13x + 4 = 0$ .

Assistant:  
② To find the sum of the reciprocals of the roots of the quadratic equation  $x^2 - 13x + 4 = 0$  [...]  
<execute>  

```
import sympy
x = sympy.Symbol('x')
roots = sympy.solve(x**2 - 13*x + 4)
print(1/roots[0] + 1/roots[1])
</execute>
```

Env:  
①  $1/(3\sqrt{17}/2 + 13/2) + 1/(13/2 - 3\sqrt{17}/2)$

Assistant:  
② The roots obtained are in symbolic form, which makes it difficult to evaluate the sum directly. Let's use the quadratic formula to explicitly find the numerical values of the roots first.  
[...]

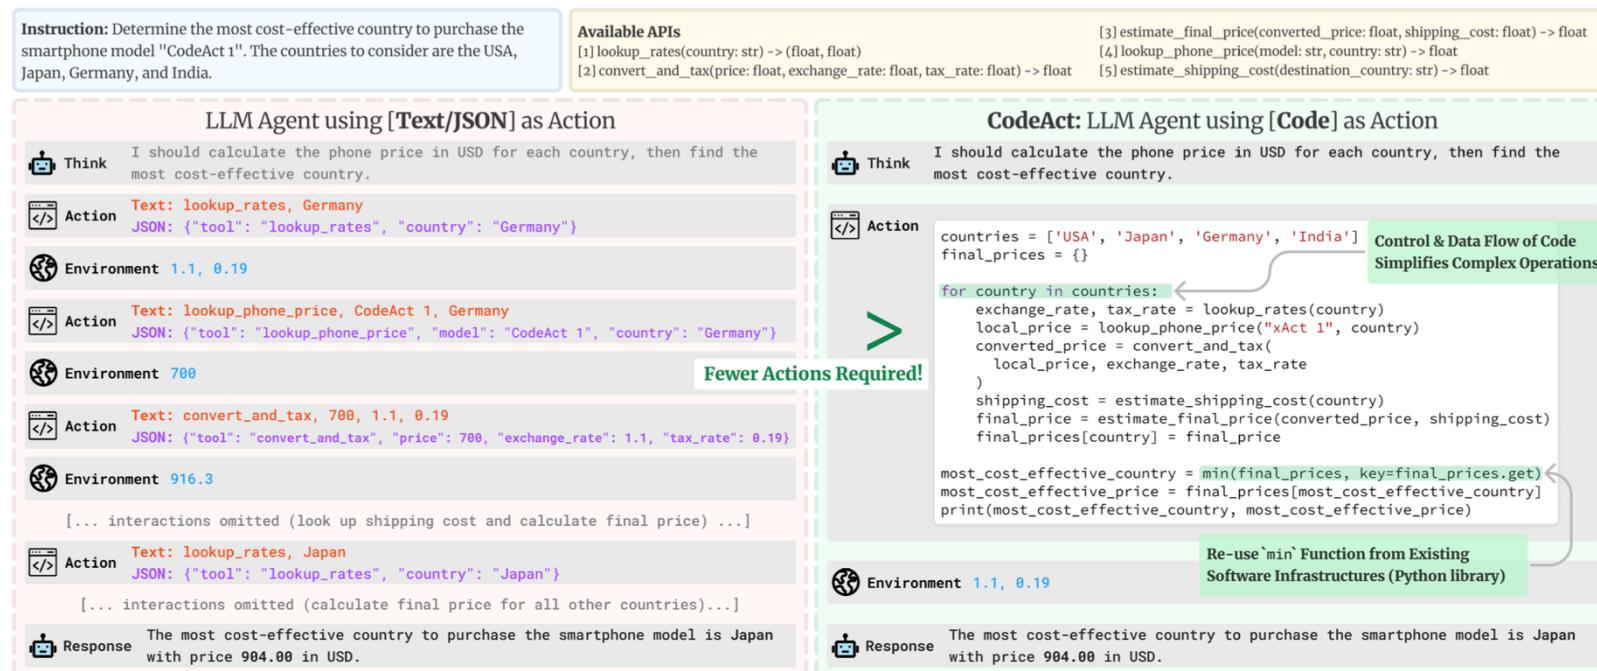


CodeACT

# Applications: Agents

Section 6.3

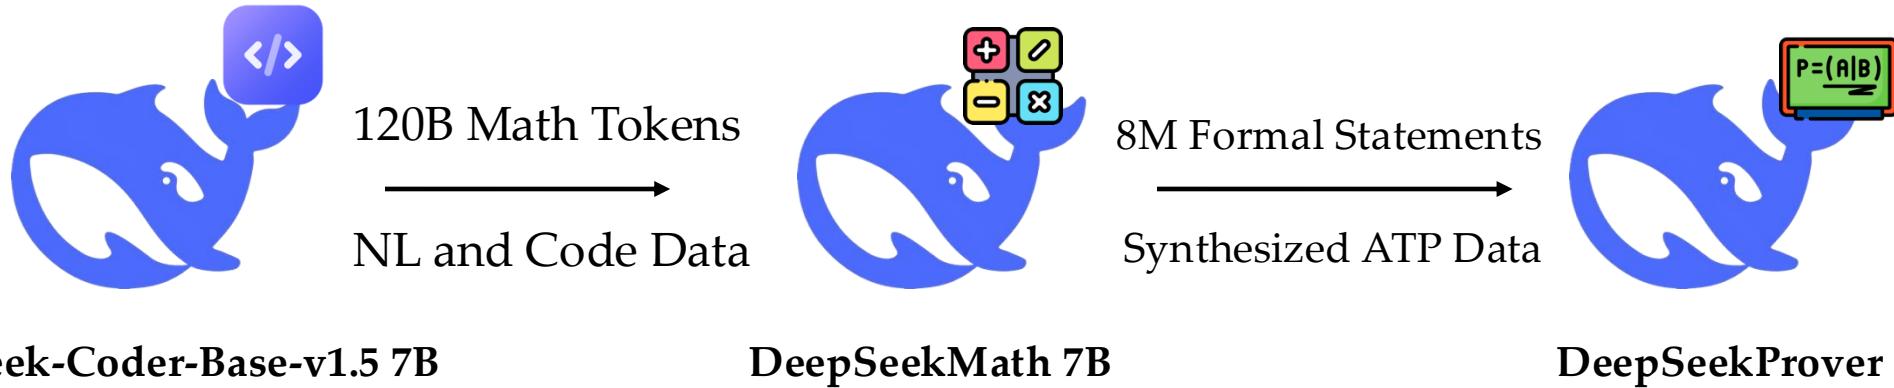
In short: Fewer actions, better efficiency



# Applications: AI4Science

Section 6.4

## Models



## Methods

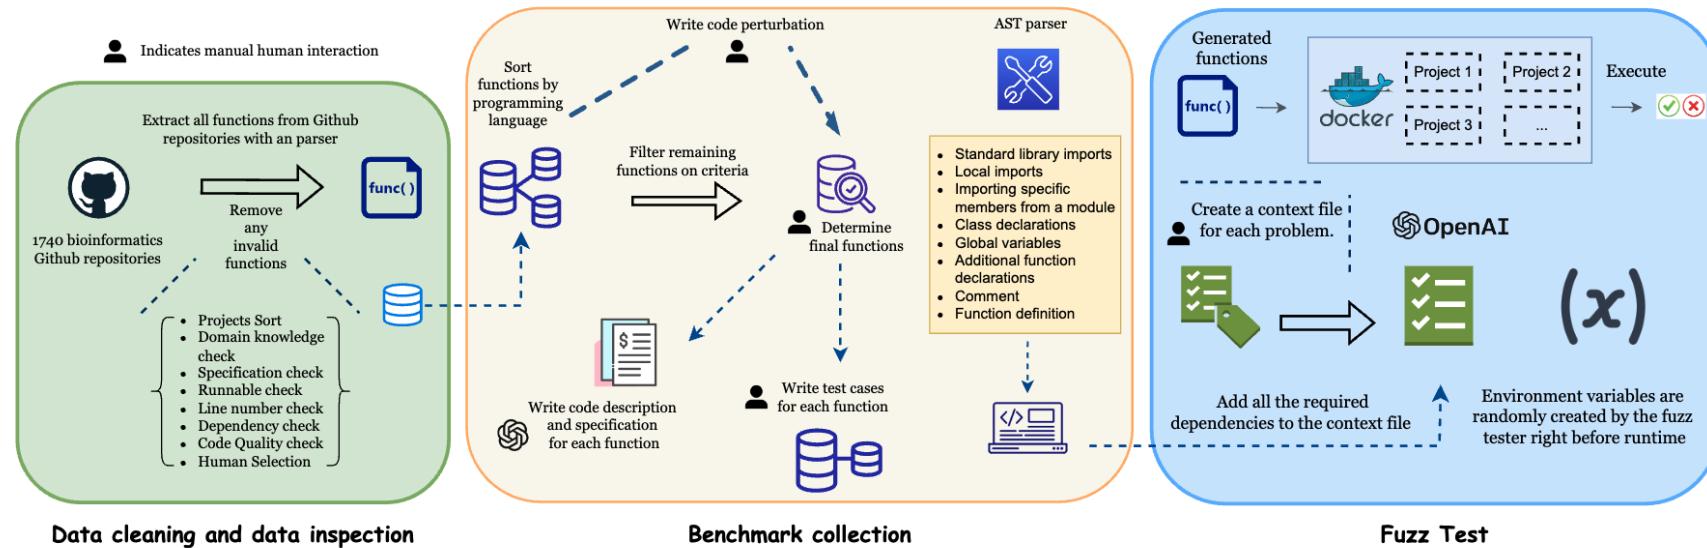
1. Biochemistry Discoveries.
2. Chemical programming languages -> automate the synthesis of chemical compounds.

# Applications: AI4Science

Section 6.4

## Biochemistry Discoveries.

Processing Bioinformatics-related GitHub repositories



# Applications: Reasoning

Program-aided language models

Coding instead of CoT

Offload “Computation” to a python interpreter

## Chain-of-Thought (Wei et al., 2022)

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 tennis balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves of bread did they have left?

Model Output

A: The bakers started with 200 loaves. They sold 93 in the morning and 39 in the afternoon. So they sold  $93 + 39 = 132$  loaves. The grocery store returned 6 loaves. So they had  $200 - 132 - 6 = 62$  loaves left.

The answer is 62.



## Program-aided Language models (this work)

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 tennis balls.  
`tennis_balls = 5`  
2 cans of 3 tennis balls each is  
`bought_balls = 2 * 3`  
tennis balls. The answer is  
`answer = tennis_balls + bought_balls`

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves of bread did they have left?

Model Output

A: The bakers started with 200 loaves  
`loaves_baked = 200`  
They sold 93 in the morning and 39 in the afternoon  
`loaves_sold_morning = 93`  
`loaves_sold_afternoon = 39`  
The grocery store returned 6 loaves.  
`loaves_returned = 6`  
The answer is  
`answer = loaves_baked - loaves_sold_morning`  
`- loaves_sold_afternoon + loaves_returned`

`>>> print(answer)`  
74



# Applications: Reasoning

Lot of techniques can be

1. Directly transferred to this scenario
2. Combine together

LLMs like DeepSeekMath-Instruct leverage such **math x code** data in training

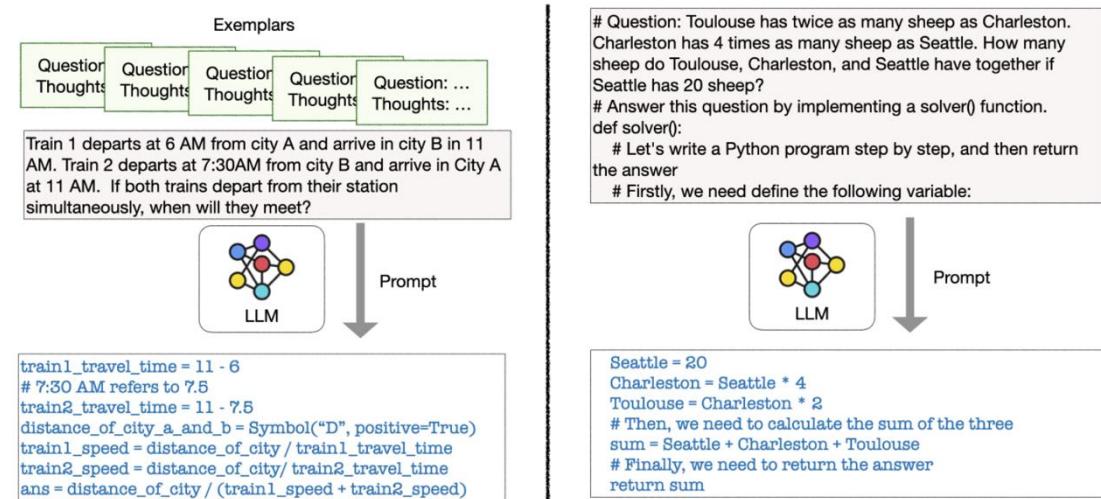
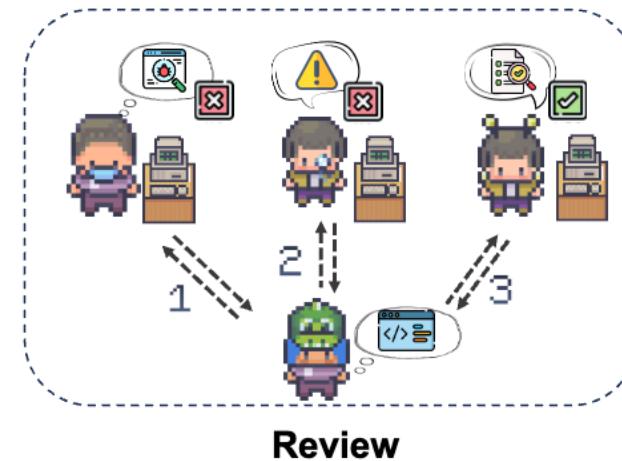


Figure 3: Left: Few-shot PoT prompting, Right: Zero-shot PoT prompting.



# Resources

## 🔗 Paper Collections / Tutorials

- [Language Models for Code](#) 
- [Evaluations and Benchmarks](#) 
- [Preference Optimization](#) 
- [Code Repair](#) 
- [Reasoning with Code Synthesis](#) 
- [Data Science](#) 
- [Corpus containing Code Data](#) 
- [Code-Based Solutions for NLP Tasks](#) 
- [Code Empowered Agents](#) 
- [Reinforcement Learning with CodeLMs](#) 
- [Code Intelligence assists AI4Science](#) 
- [Software Development](#) 
- [Multilingual](#) 
- [Multimodal Code Generation](#) 
- [Awesome Slides, Talks and Blogs](#) 

## Recent Work on Code Intelligence (Welcome PR)

- [Codel/O: Condensing Reasoning Patterns via Code Input-Output Prediction](#) 2025.2
- [Competitive Programming with Large Reasoning Models](#) 2025.2
- [EpiCoder: Encompassing Diversity and Complexity in Code Generation](#) 2025.1
- [WarriorCoder: Learning from Expert Battles to Augment Code Large Language Models](#) 2024.12
- [FullStack Bench: Evaluating LLMs as Full Stack Coders](#) 2024.12
- [CodeDPO: Aligning Code Models with Self Generated and Verified Source Code](#) 2024.11
- [OpenCoder: The Open Cookbook for Top-Tier Code Large Language Models](#) 2024.11
- [Qwen2.5-Coder Series: Powerful, Diverse, Practical.](#) 2024.11

<https://github.com/QiushiSun/Awesome-Code-Intelligence>

The background of the slide is a solid blue color with a faint, abstract pattern of white lines and shapes resembling a circuit board or a network diagram.

# Thanks for listening

Contact: [qiushisun@connect.hku.hk](mailto:qiushisun@connect.hku.hk)